



Image by Connie Zhou

Failures

A 2000-machine cluster will have
1 to 10 machine failures per day.
This is not a problem: it's normal.

Kubernetes


Greek for “*Helmsman*”; also the root of the words “*governor*” and “*cybernetic*”

- Manages container clusters
- Inspired and informed by Google’s experiences and internal systems
- Supports multiple cloud and bare-metal environments
- Supports multiple container runtimes
- **100% Open source**, written in Go

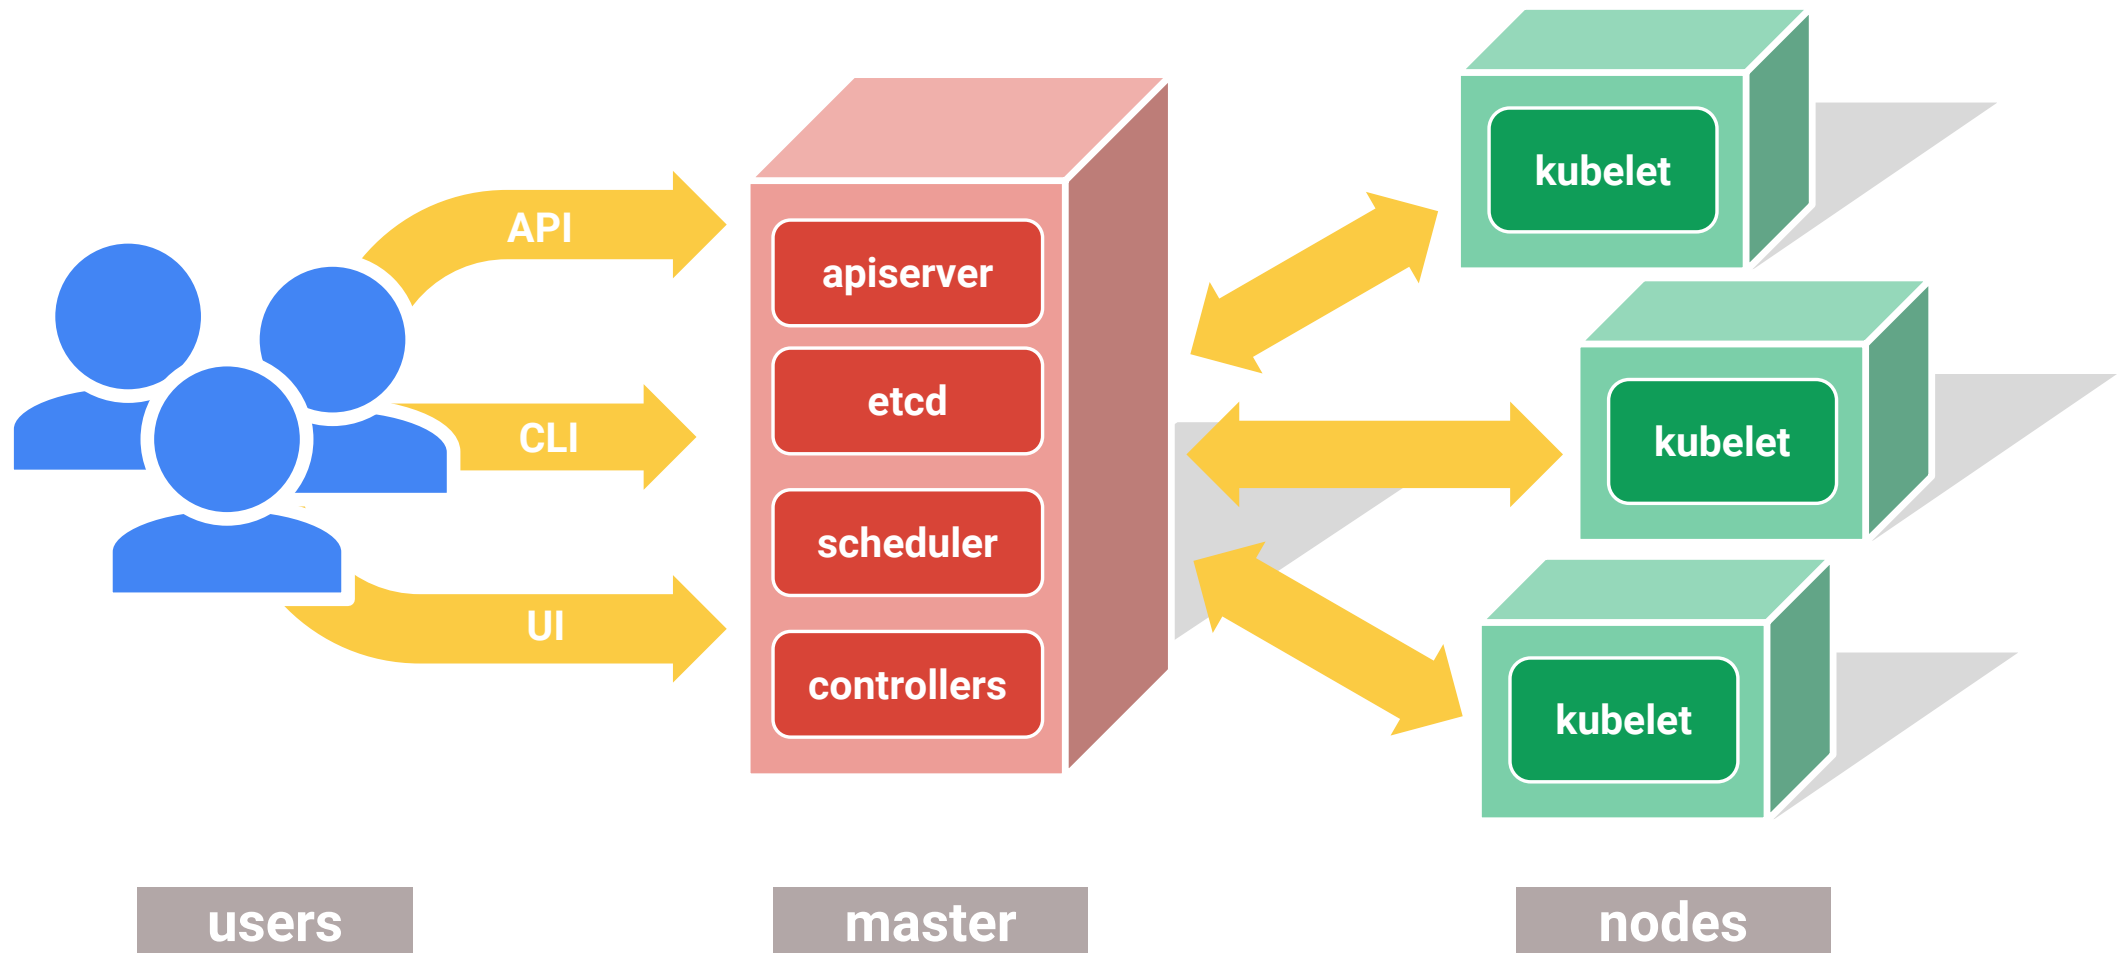
Manage applications, not machines



All you really care about



The 10000 foot view



Container clusters: A story in two parts



Container clusters: A story in two parts

1. Setting up the cluster

- Choose a cloud: GCE, AWS, Azure, Rackspace, on-premises, ...
- Choose a node OS: CoreOS, Atomic, RHEL, Debian, CentOS, Ubuntu, ...
- Provision machines: Boot VMs, install and run kube components, ...
- Configure networking: IP ranges for Pods, Services, SDN, ...
- Start cluster services: DNS, logging, monitoring, ...
- Manage nodes: kernel upgrades, OS updates, hardware failures...

Not the easy or fun part, but unavoidable

This is where things like **Google Container Engine (GKE)** really help



Container clusters: A story in two parts

2. Using the cluster

- Run Pods & Containers
- ReplicaSets & Deployments & DaemonSets & StatefulSets
- Services & Volumes & Secrets & Autoscalers

This is the fun part!

A distinct set of problems from cluster setup and management

Don't make developers deal with cluster administration!

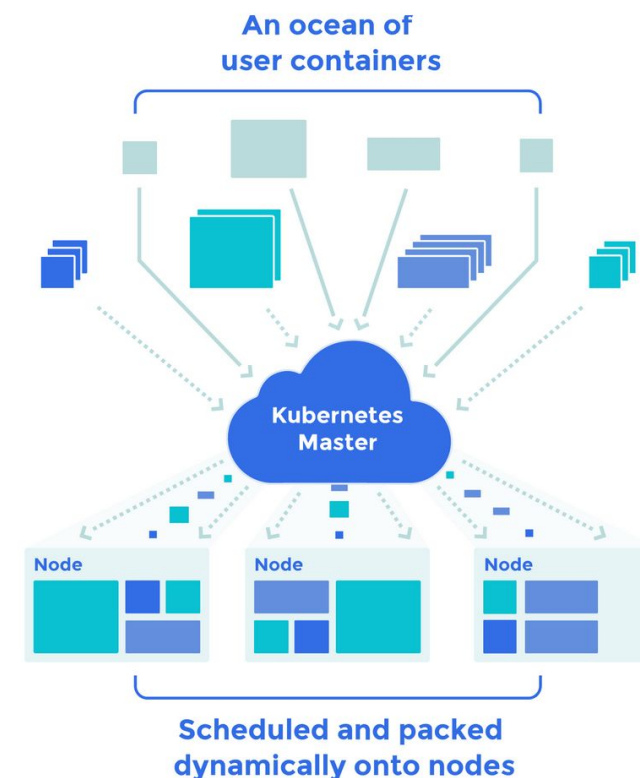
Accelerate development by focusing on the applications, not the cluster



Kubernetes: a Cloud OS?

Perhaps grandiose, but attempts at “Cloud OS” primitives:

- **Scheduling:** Decide where my containers should run
- **Lifecycle and health:** Keep my containers running despite failures
- **Scaling:** Make sets of containers bigger or smaller
- **Naming and discovery:** Find where my containers are now
- **Load balancing:** Distribute traffic across a set of containers
- **Storage volumes:** Provide data to containers
- **Logging and monitoring:** Track what’s happening with my containers
- **Debugging and introspection:** Enter or attach to containers
- **Identity and authorization:** Control who can do things to my containers



Workload Portability




Workload portability

Goal: Avoid vendor lock-in

Runs in many environments, including “bare metal” and “your laptop”

The API and the implementation are 100% open

The whole system is modular and replaceable



Workload portability


Goal: Write once, run anywhere^{*}

Don't force apps to know about concepts that are cloud-provider-specific

Examples of this:

- Network model
- Ingress
- Service load-balancers
- PersistentVolumes

** approximately*




Workload portability

Result: Portability

Build your apps on-prem, lift-and-shift into cloud when you are ready

Don't get stuck with a platform that doesn't work for you


Put your app on wheels and move it whenever and wherever you need



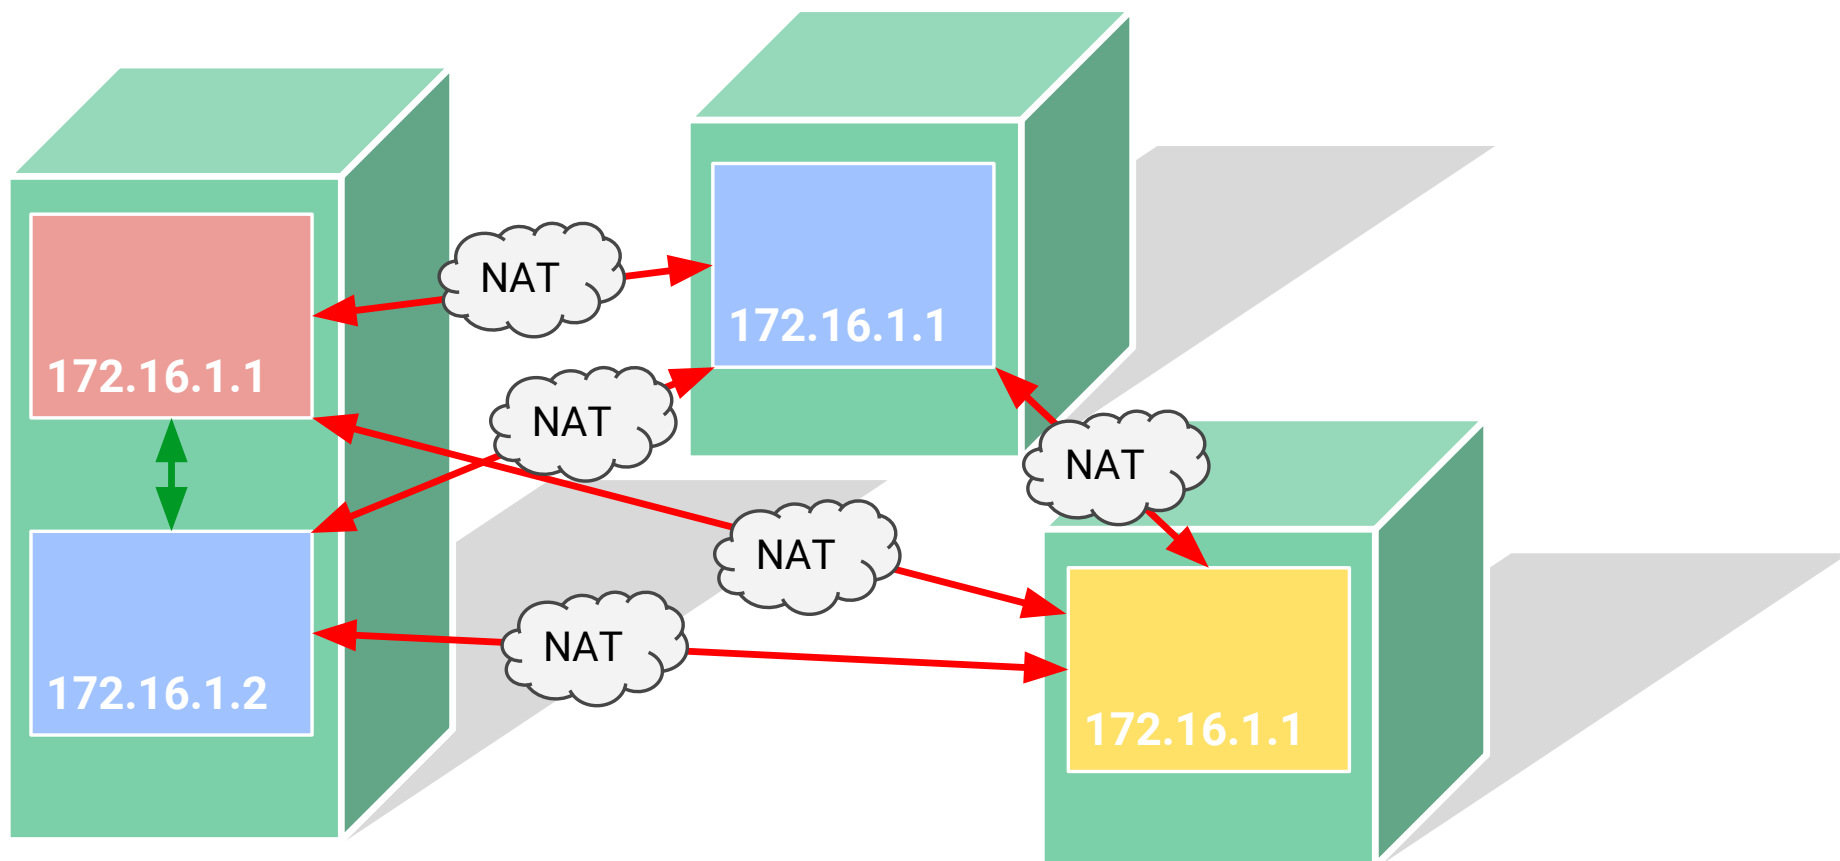
Networking




Docker networking




Docker networking



Port mapping



Port mapping



Kubernetes networking

IPs are **cluster-scoped**

- vs docker default private IP

Pods can reach each other directly


- even across nodes

No **brokering** of port numbers


- too complex, why bother?

This is a **fundamental requirement**

- can be L3 routed
- can be underlayed (cloud)
- can be overlayed (SDN)



Kubernetes networking



Pods



Pods

Small group of containers & volumes

Tightly coupled

The atom of scheduling & placement

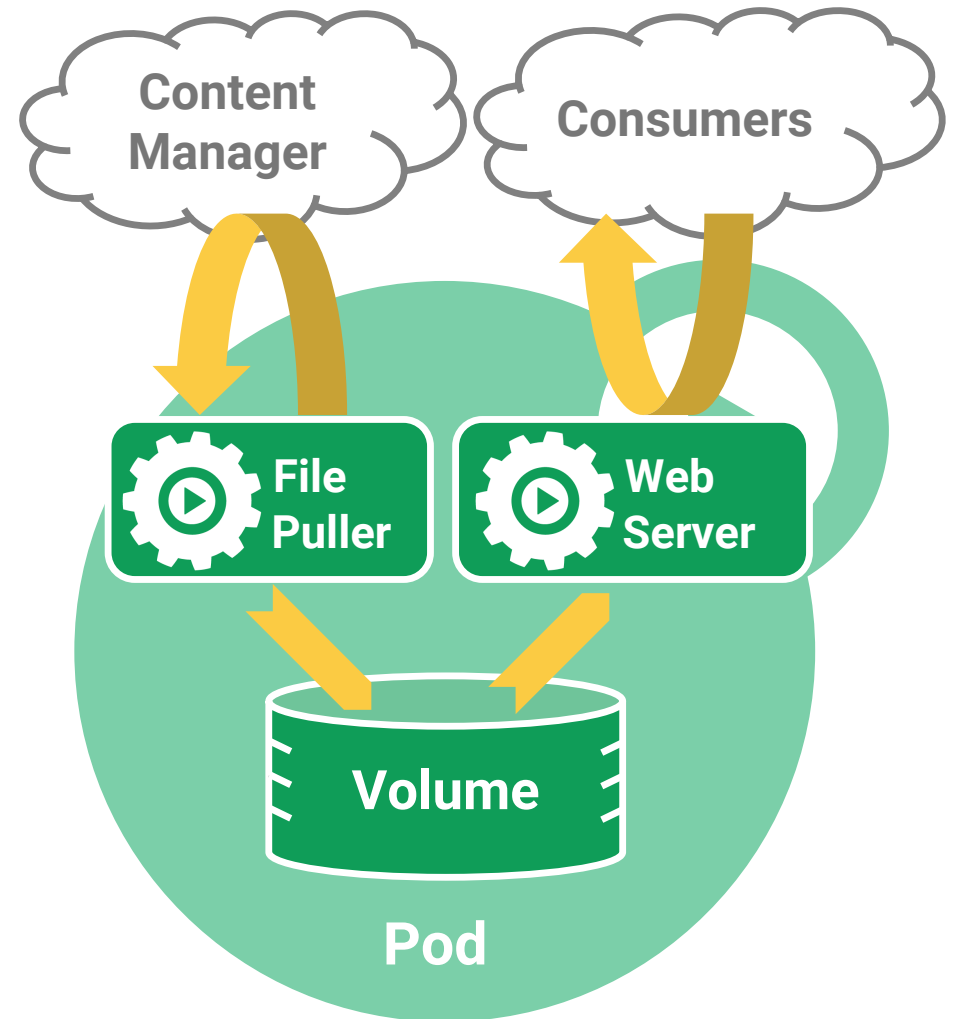
Shared namespace

- share IP address & localhost
- share IPC, etc.

Managed lifecycle

- bound to a node, restart in place
- can die, cannot be reborn with same ID

Example: data puller & web server



Volumes

Pod-scoped storage

Support many types of volume plugins

- Empty dir (and tmpfs)
- Host path
- Git repository
- GCE Persistent Disk
- AWS Elastic Block Store
- Azure File Storage
- iSCSI
- Flocker
- NFS
- vSphere
- GlusterFS
- Ceph File and RBD
- Cinder
- FibreChannel
- Secret, ConfigMap, DownwardAPI
- Flex (exec a binary)
- ...



Labels & Selectors



Labels

Arbitrary metadata

Attached to any API object

Generally represent **identity**

Queryable by **selectors**


- think SQL *'select ... where ...'*

The **only** grouping mechanism


- pods under a ReplicaSet
- pods in a Service
- capabilities of a node (constraints)



Selectors



Selectors



App = MyApp



Selectors




App = MyApp, Role = FE

Selectors




App = MyApp, Role = BE

Selectors



App = MyApp, Phase = prod

Selectors



App = MyApp, Phase = test

Replication



ReplicaSets

A simple control loop

Runs out-of-process wrt API server

One job: ensure N copies of a pod

- grouped by a selector
- too few? start some
- too many? kill some


Layered on top of the public Pod API

Replicated pods are **fungible**

- No implied order or identity

ReplicaSet

- name = "my-rc"
- selector = {"App": "MyApp"}
- template = { ... }
- replicas = 4



Control loops: the Reconciler Pattern

Drive **current state** -> **desired state**


Act independently

APIs - **no shortcuts** or back doors

Observed state is truth*

Recurring pattern in the system

Example: ReplicaSet



* Observations are really stale caches of what once was your view of truth.

Services



Services

A group of **pods that work together**

- grouped by a selector

Defines access policy

- “load balanced” or “headless”


Can have a stable **virtual IP** and port

- also a DNS name


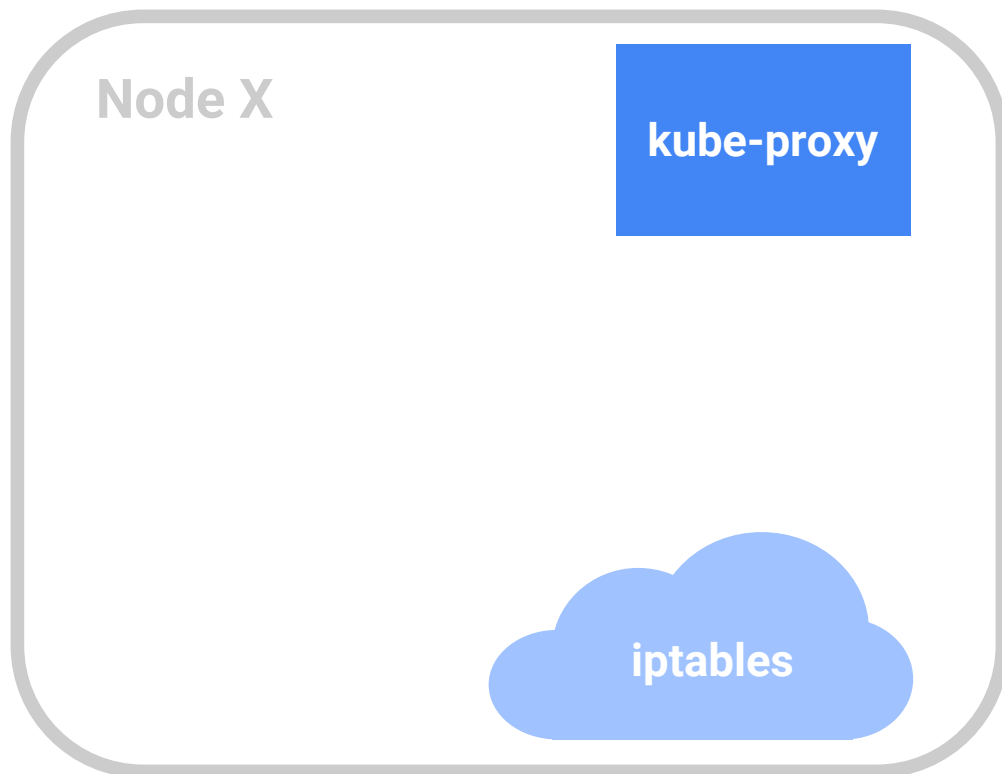
VIP is managed by *kube-proxy*

- watches all services
- updates iptables when backends change
- default implementation - can be replaced!

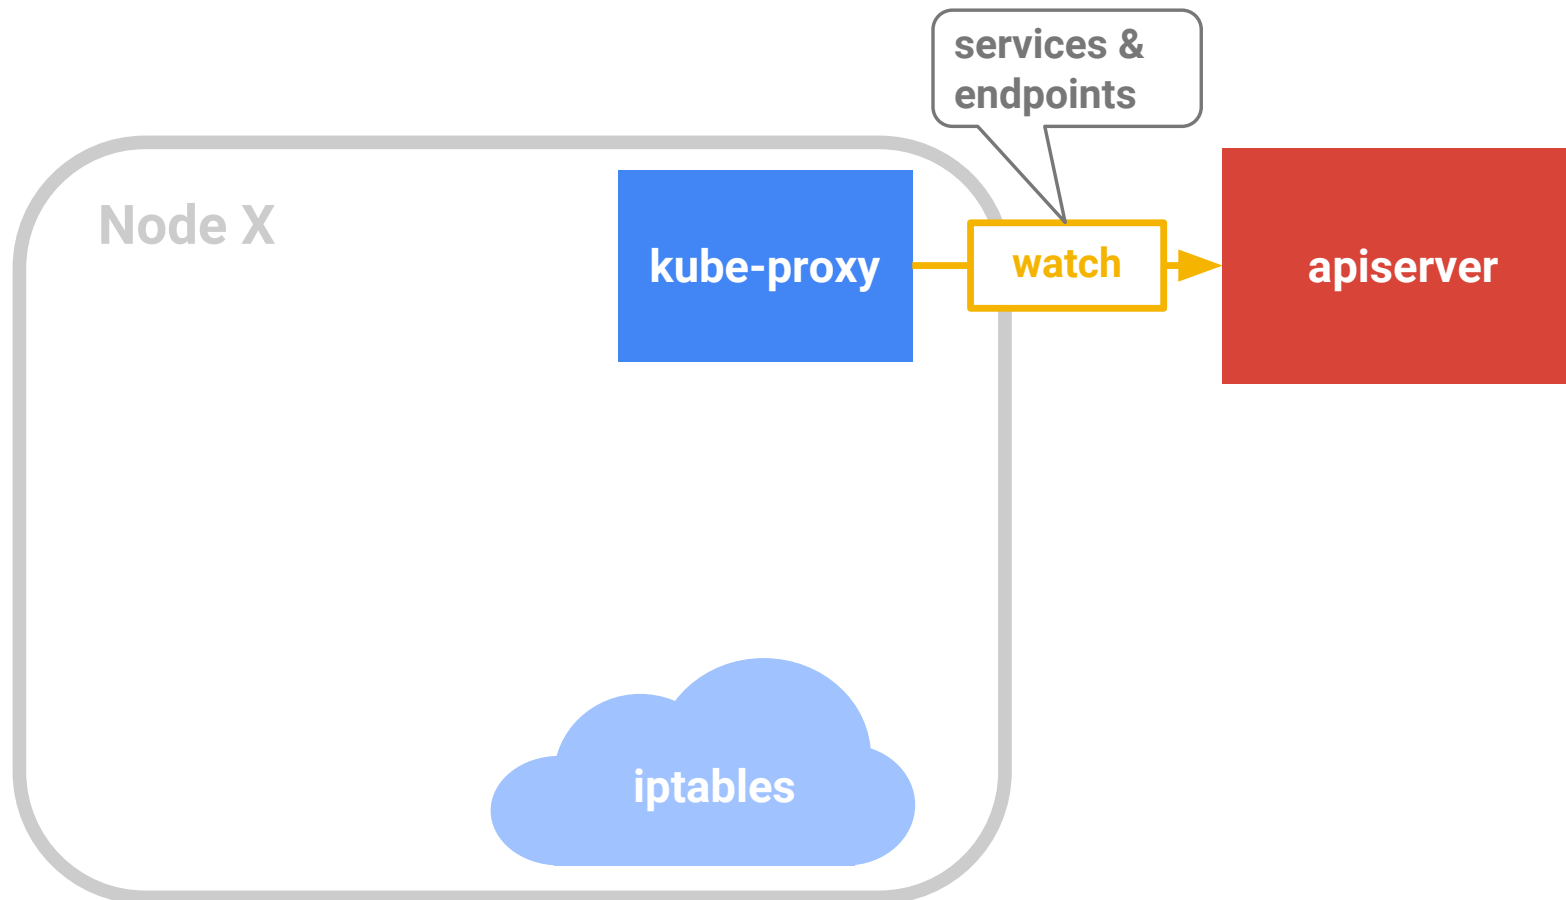
Hides complexity



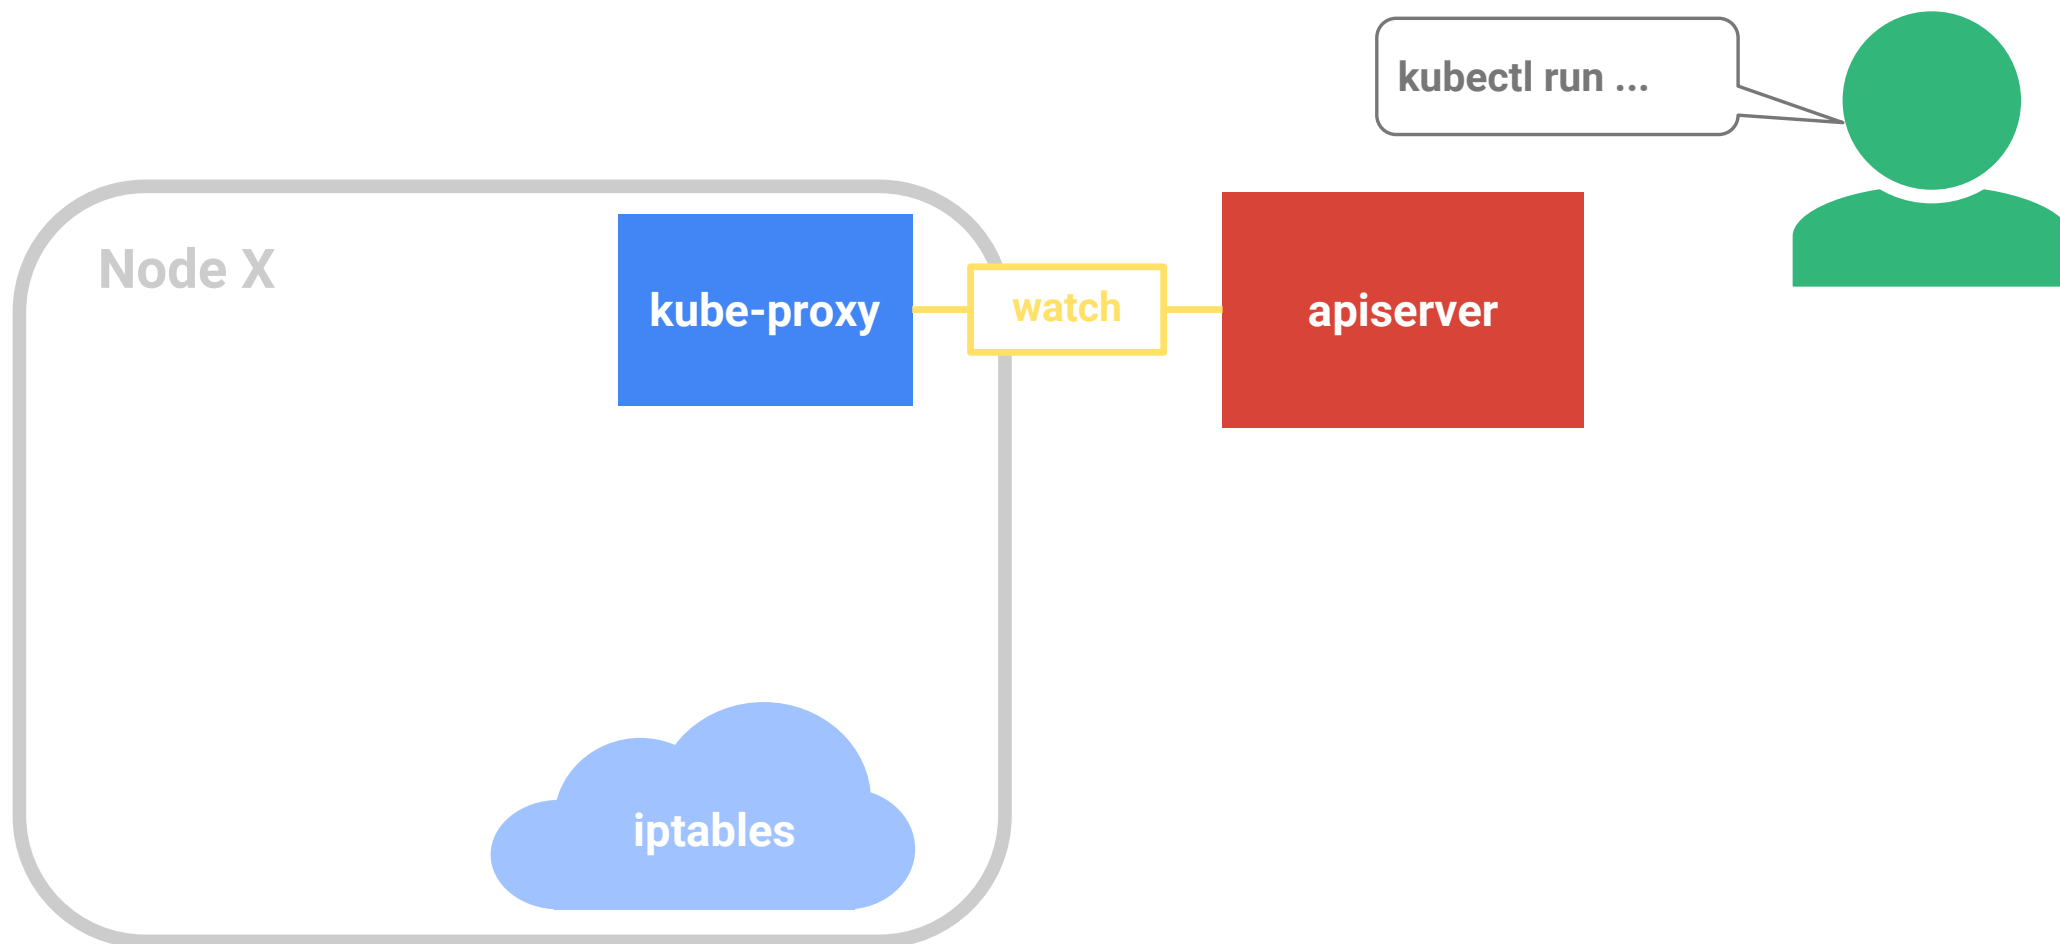
iptables kube-proxy




iptables kube-proxy



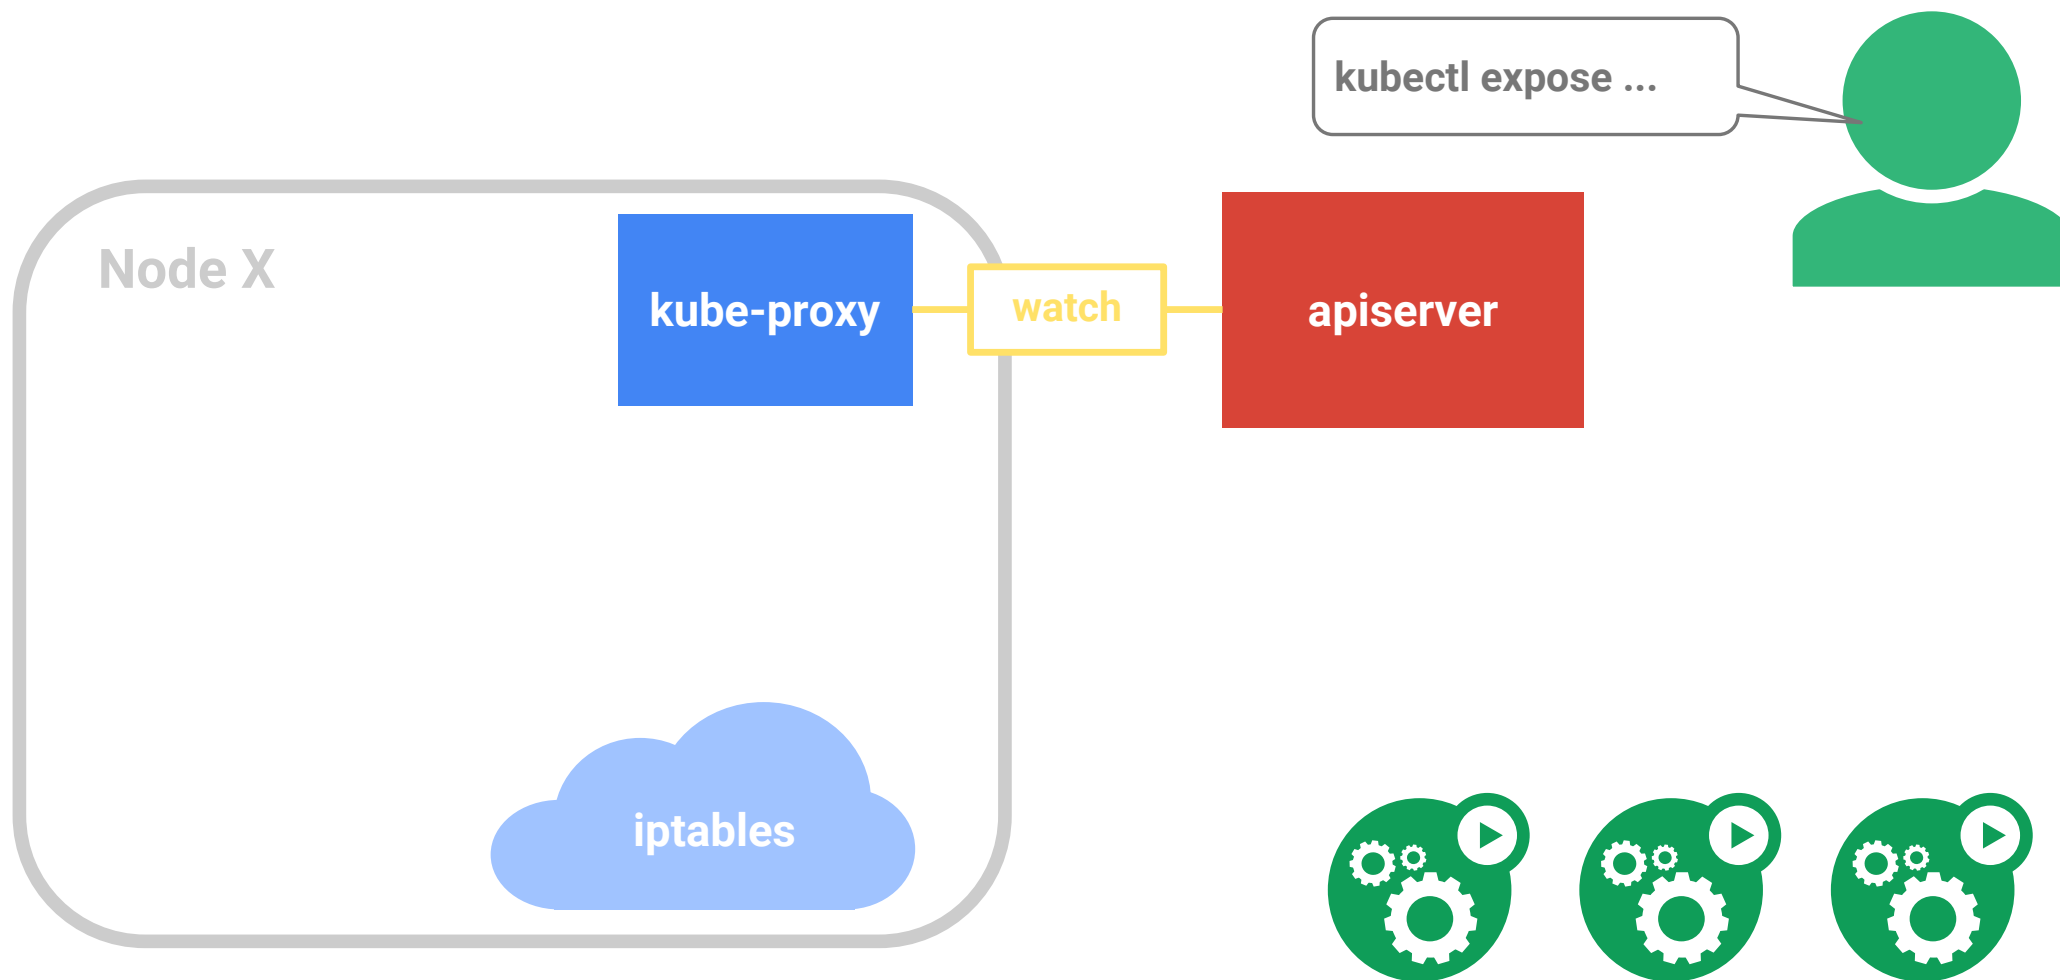
iptables kube-proxy



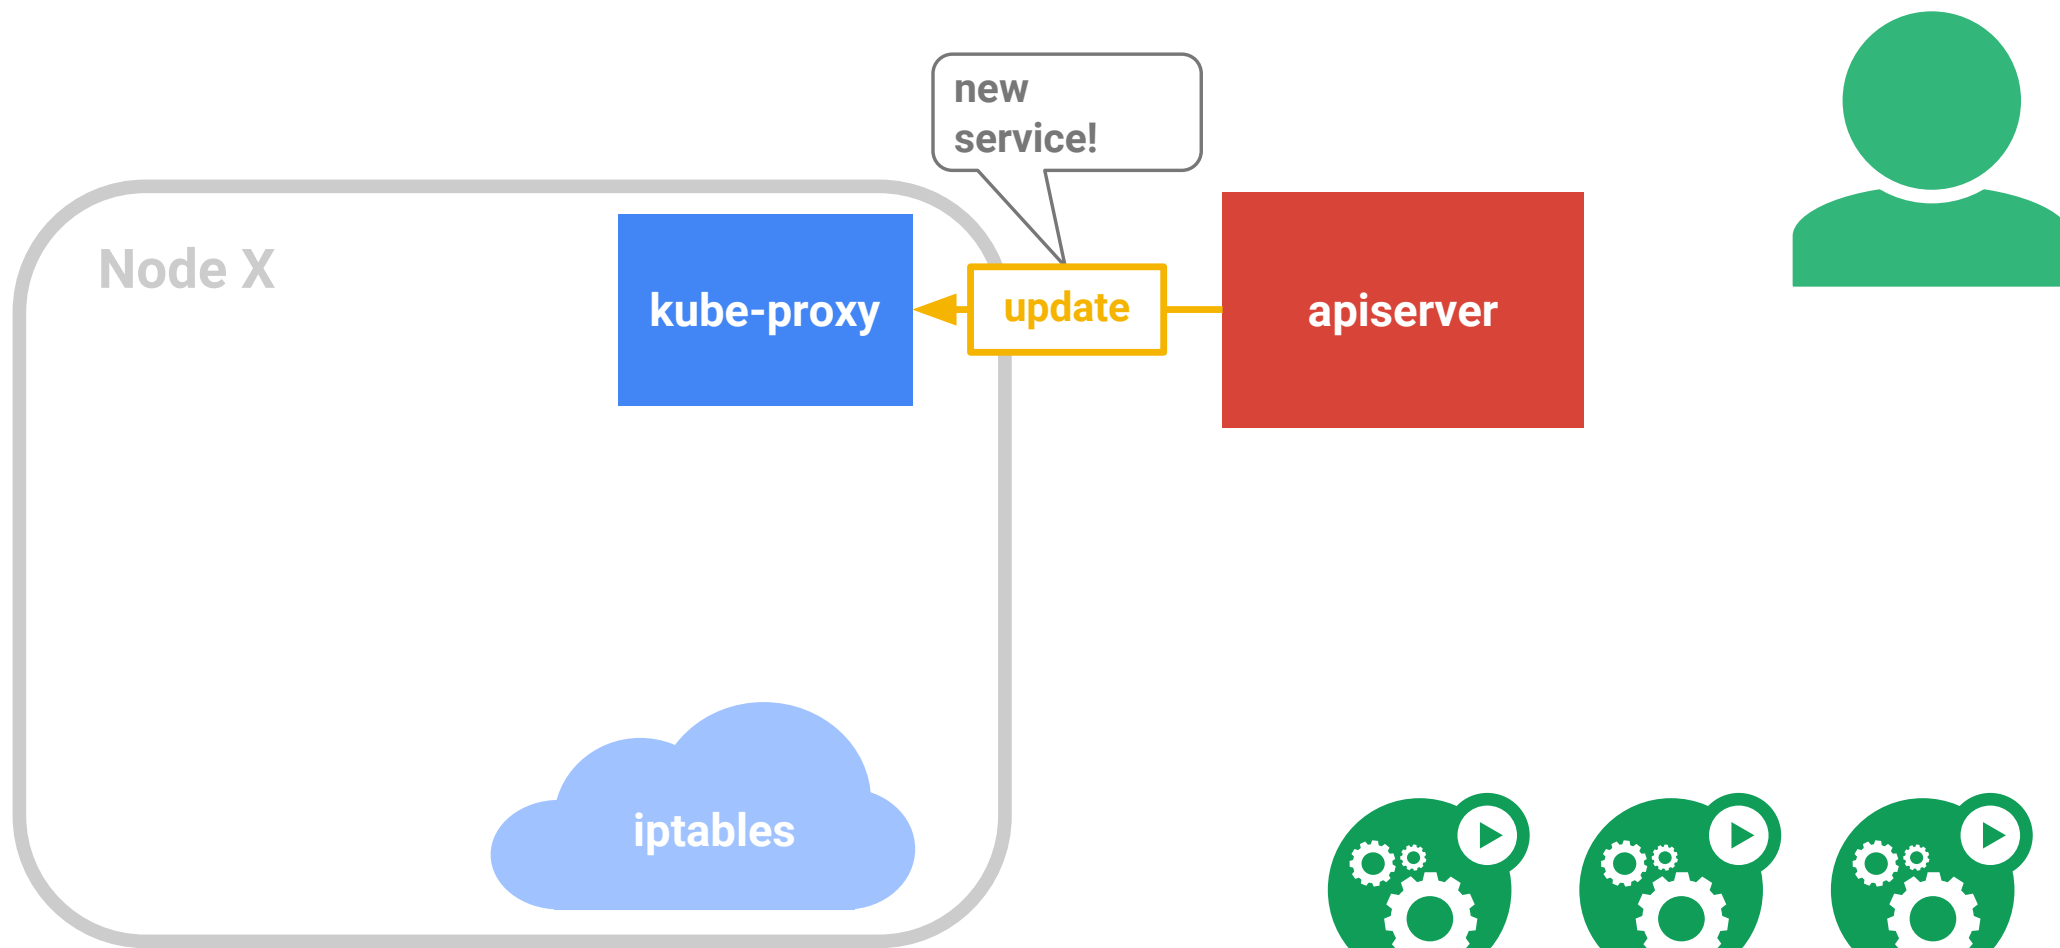
iptables kube-proxy




iptables kube-proxy




iptables kube-proxy



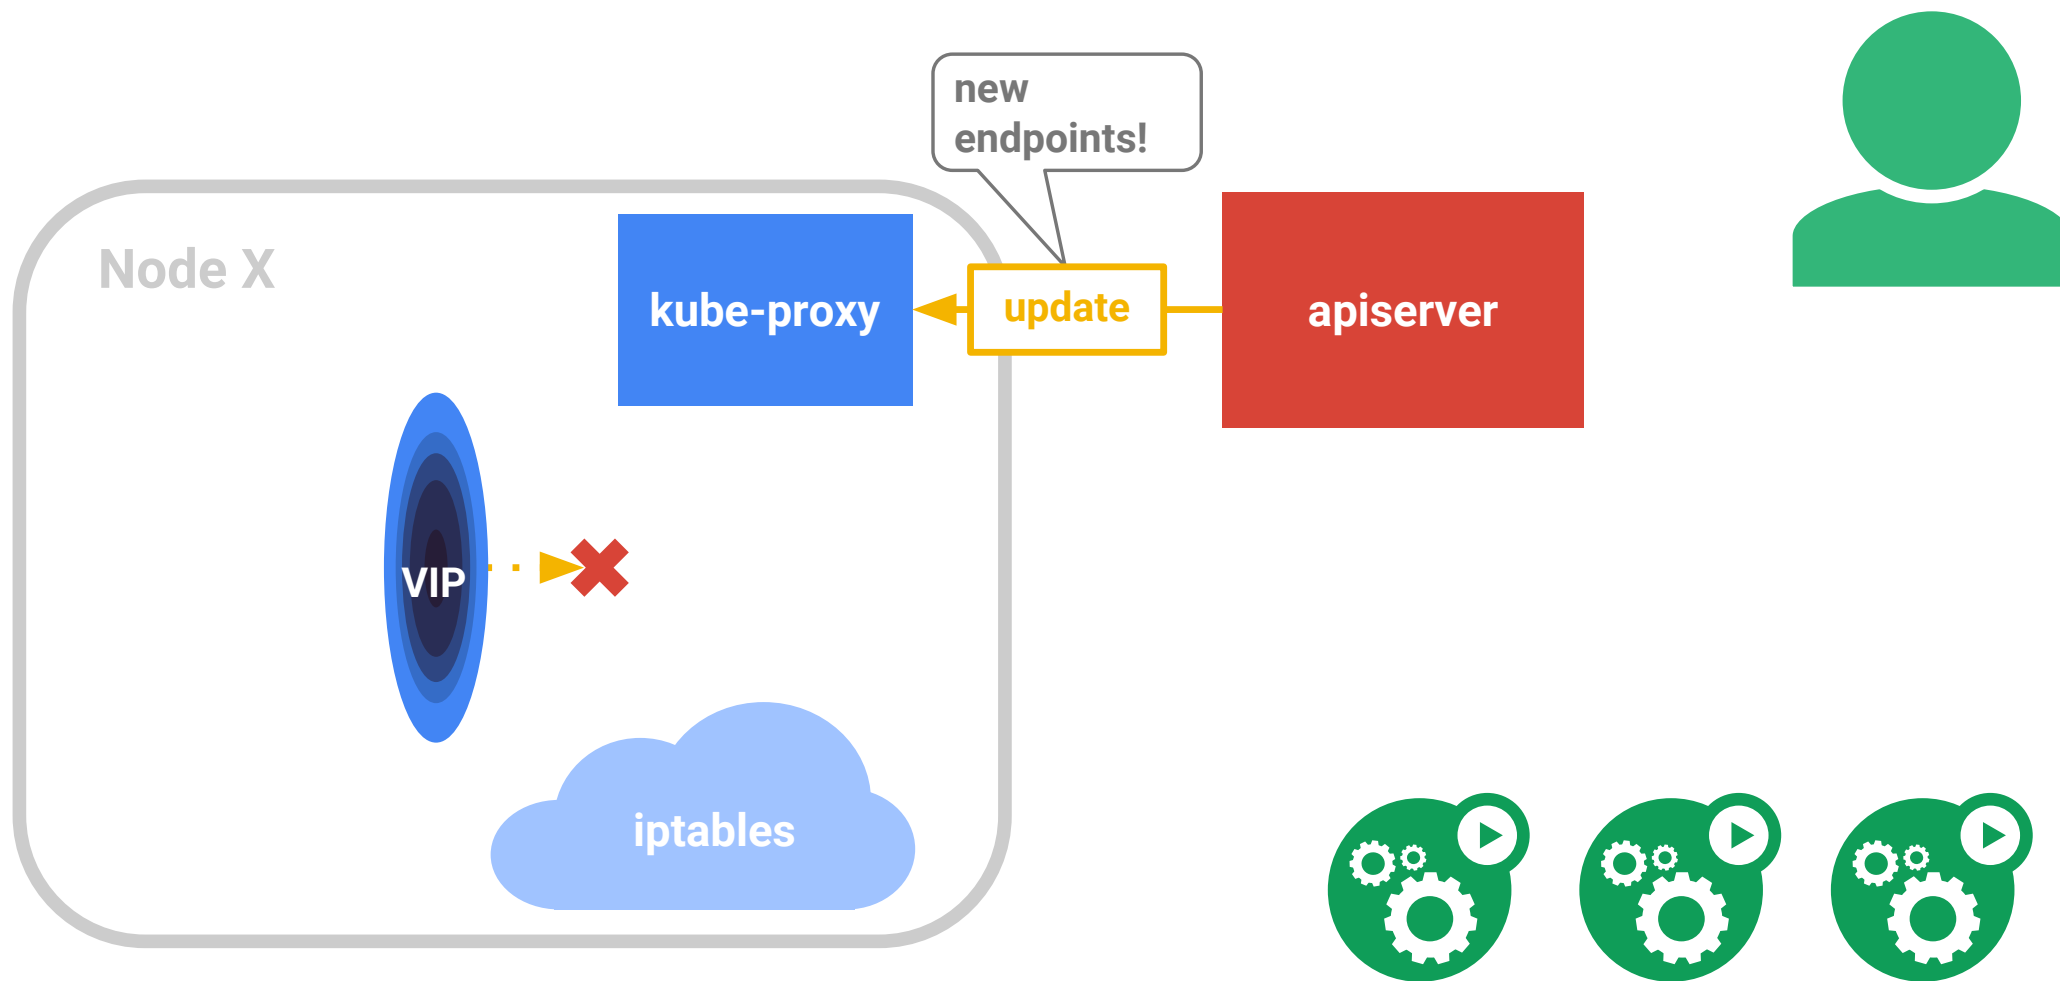
iptables kube-proxy




iptables kube-proxy




iptables kube-proxy



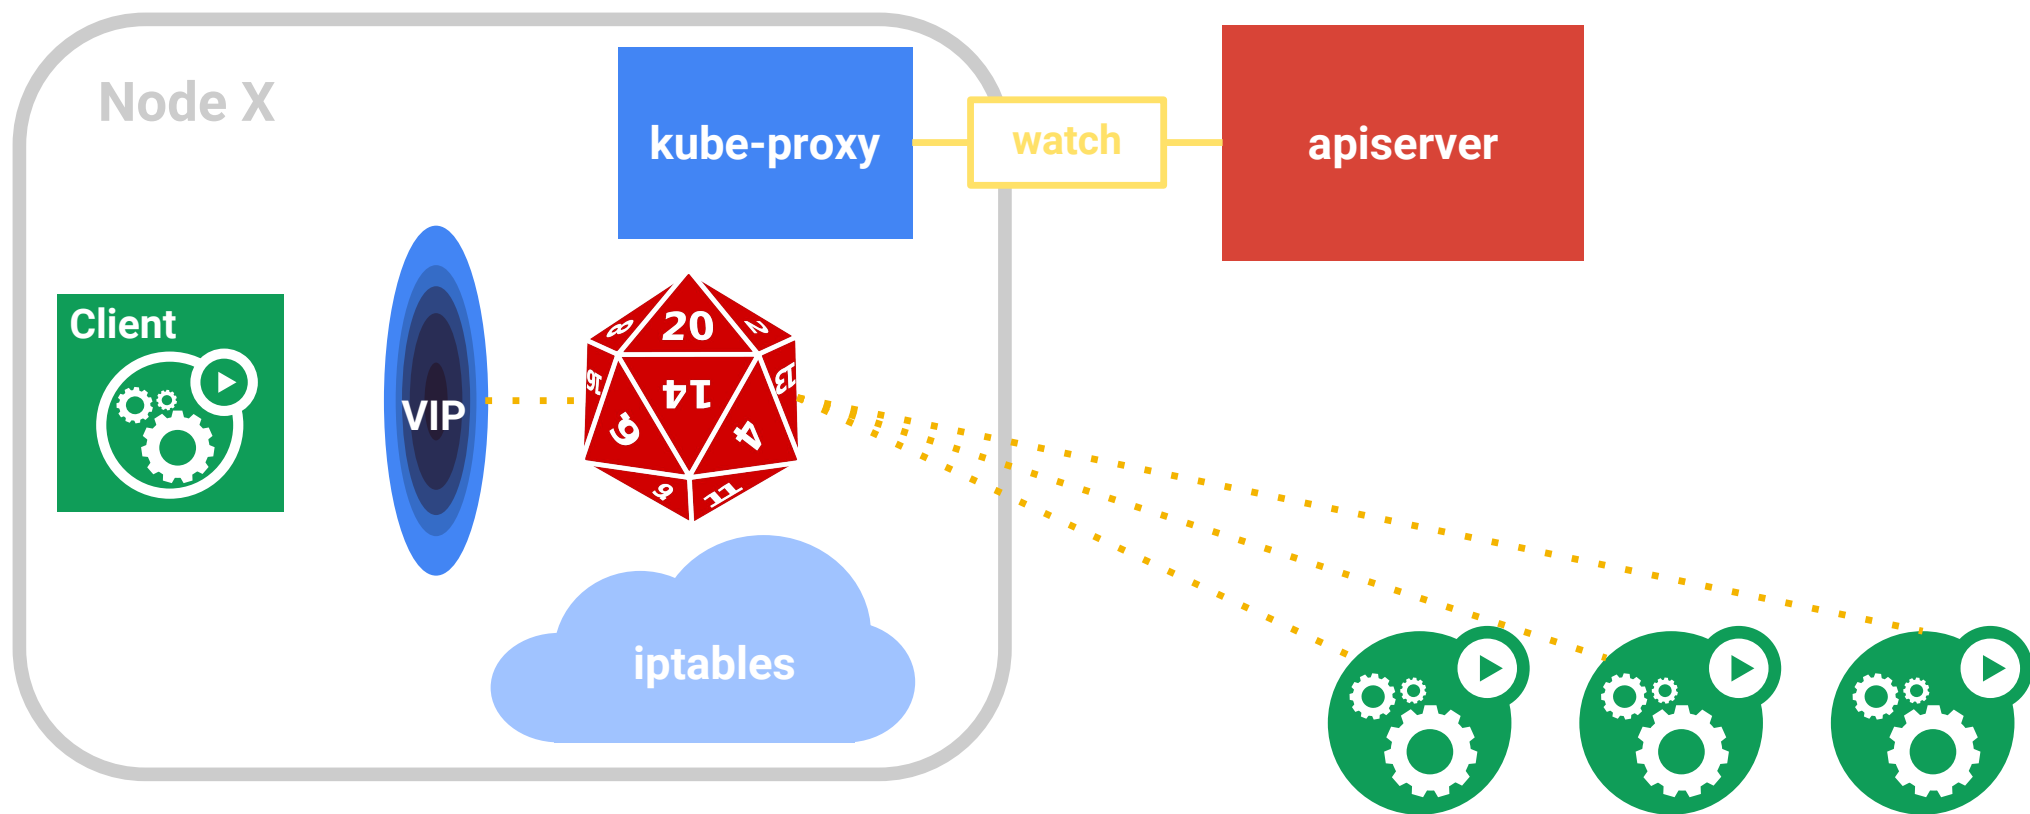
iptables kube-proxy



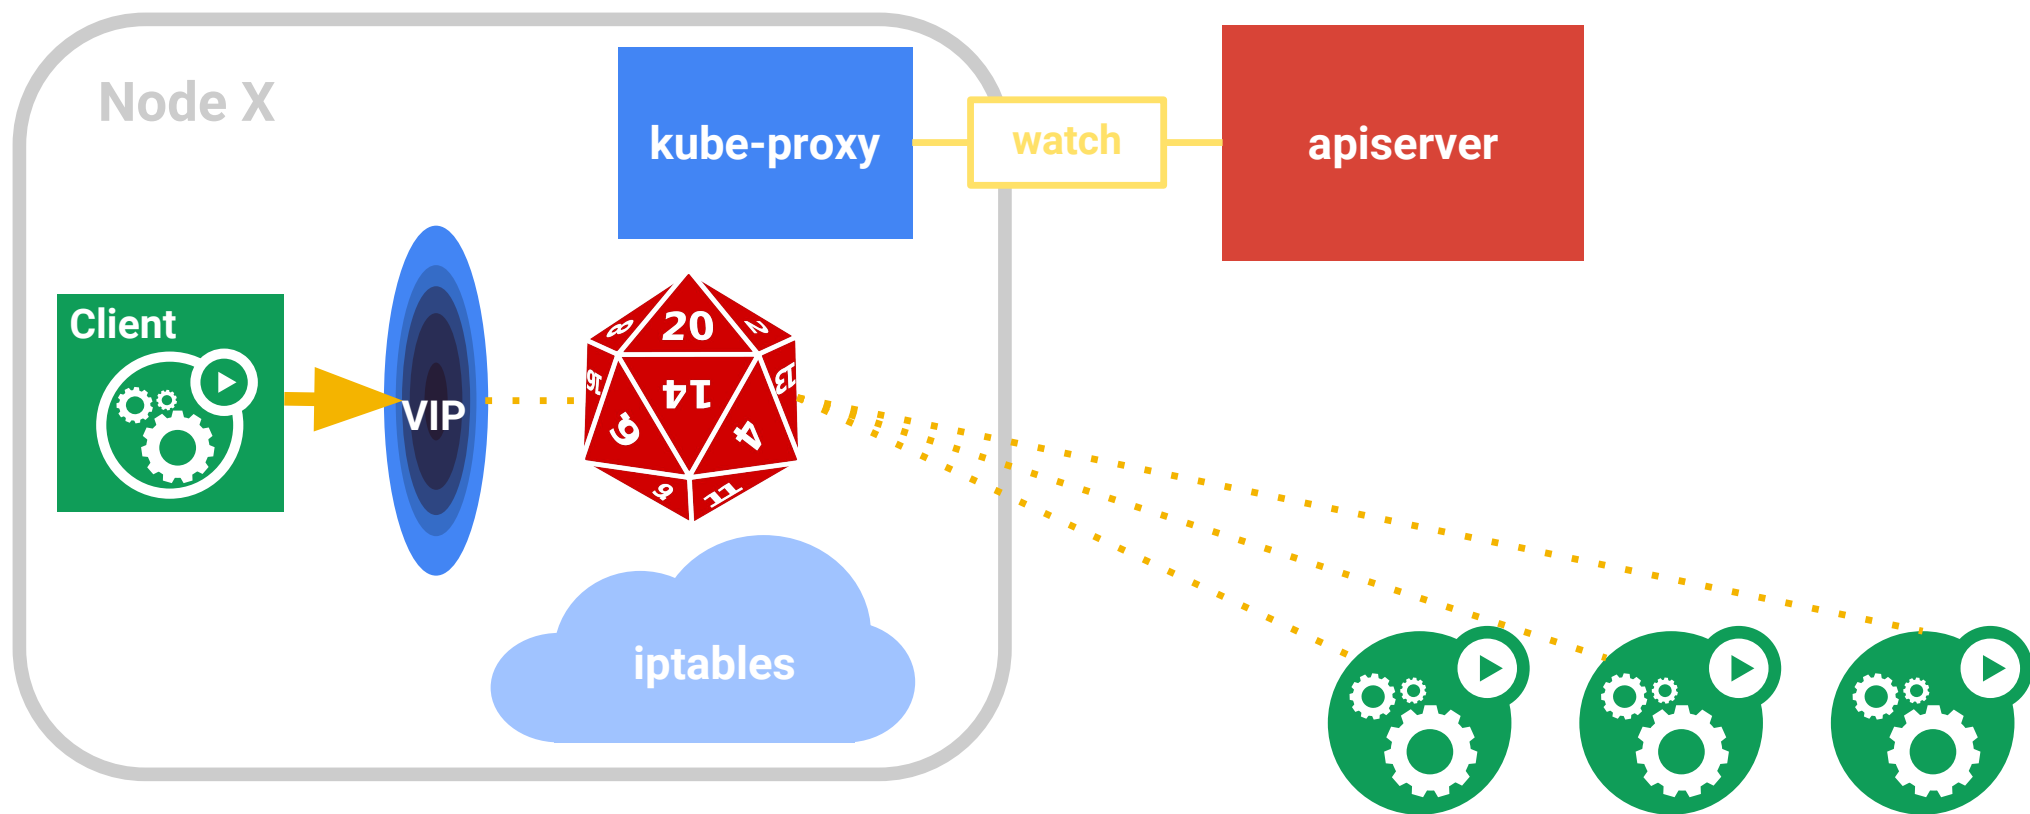
iptables kube-proxy



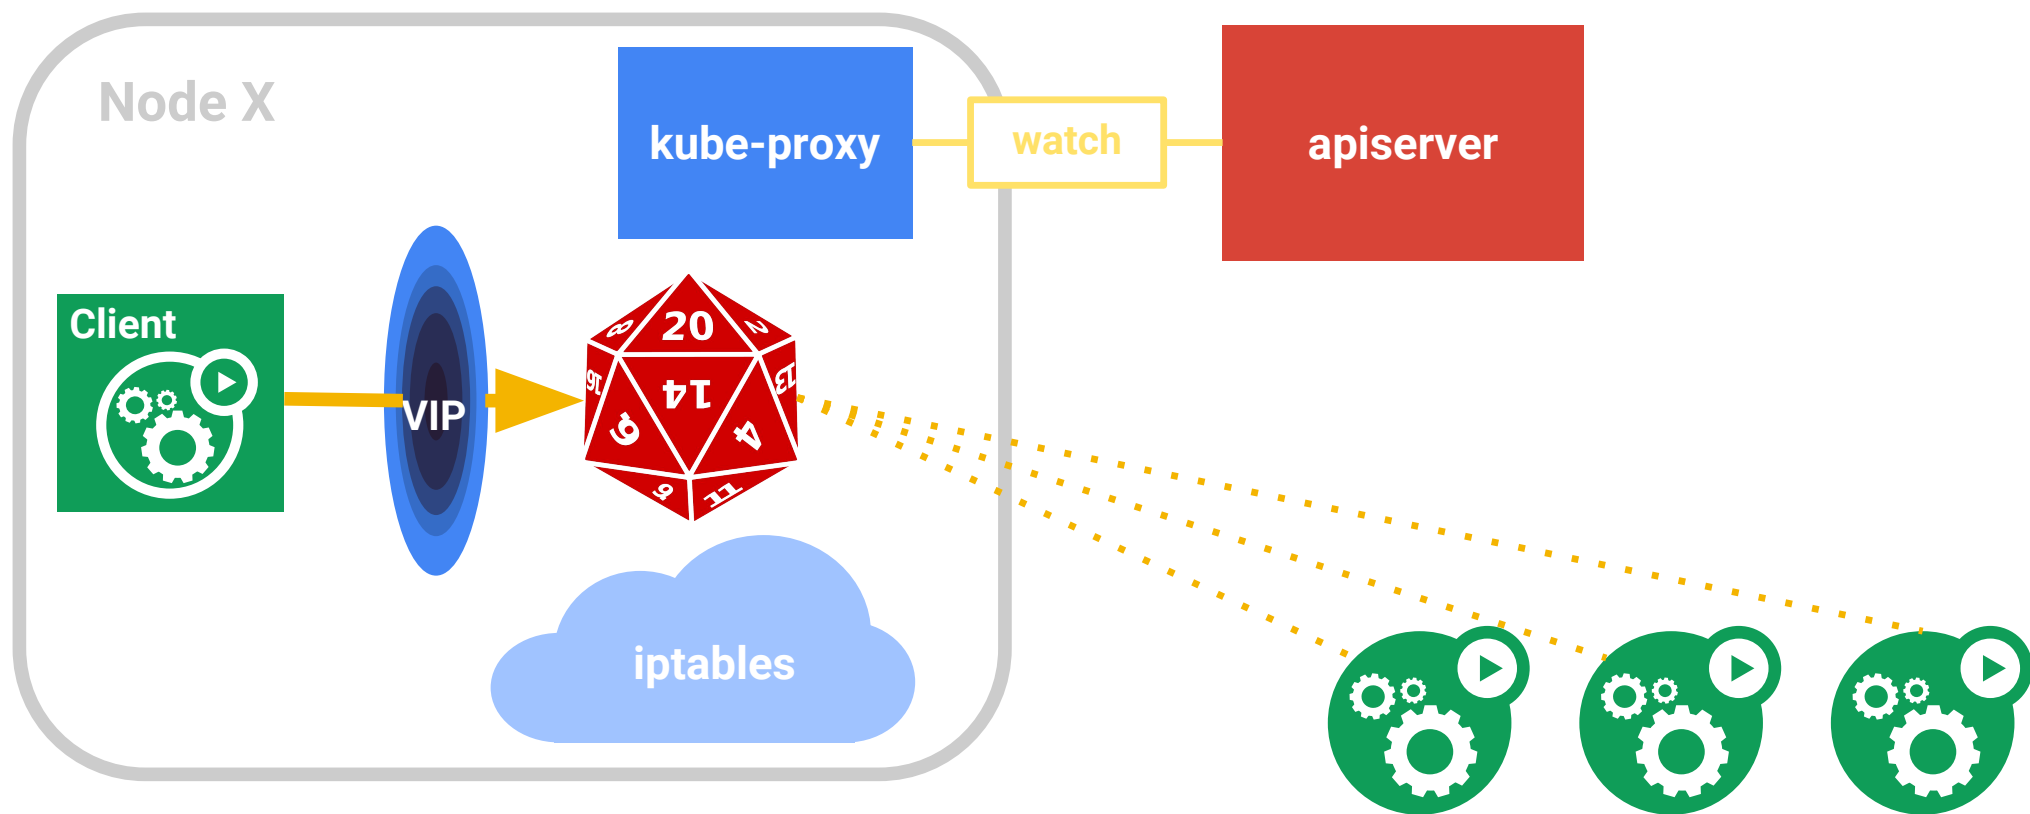
iptables kube-proxy



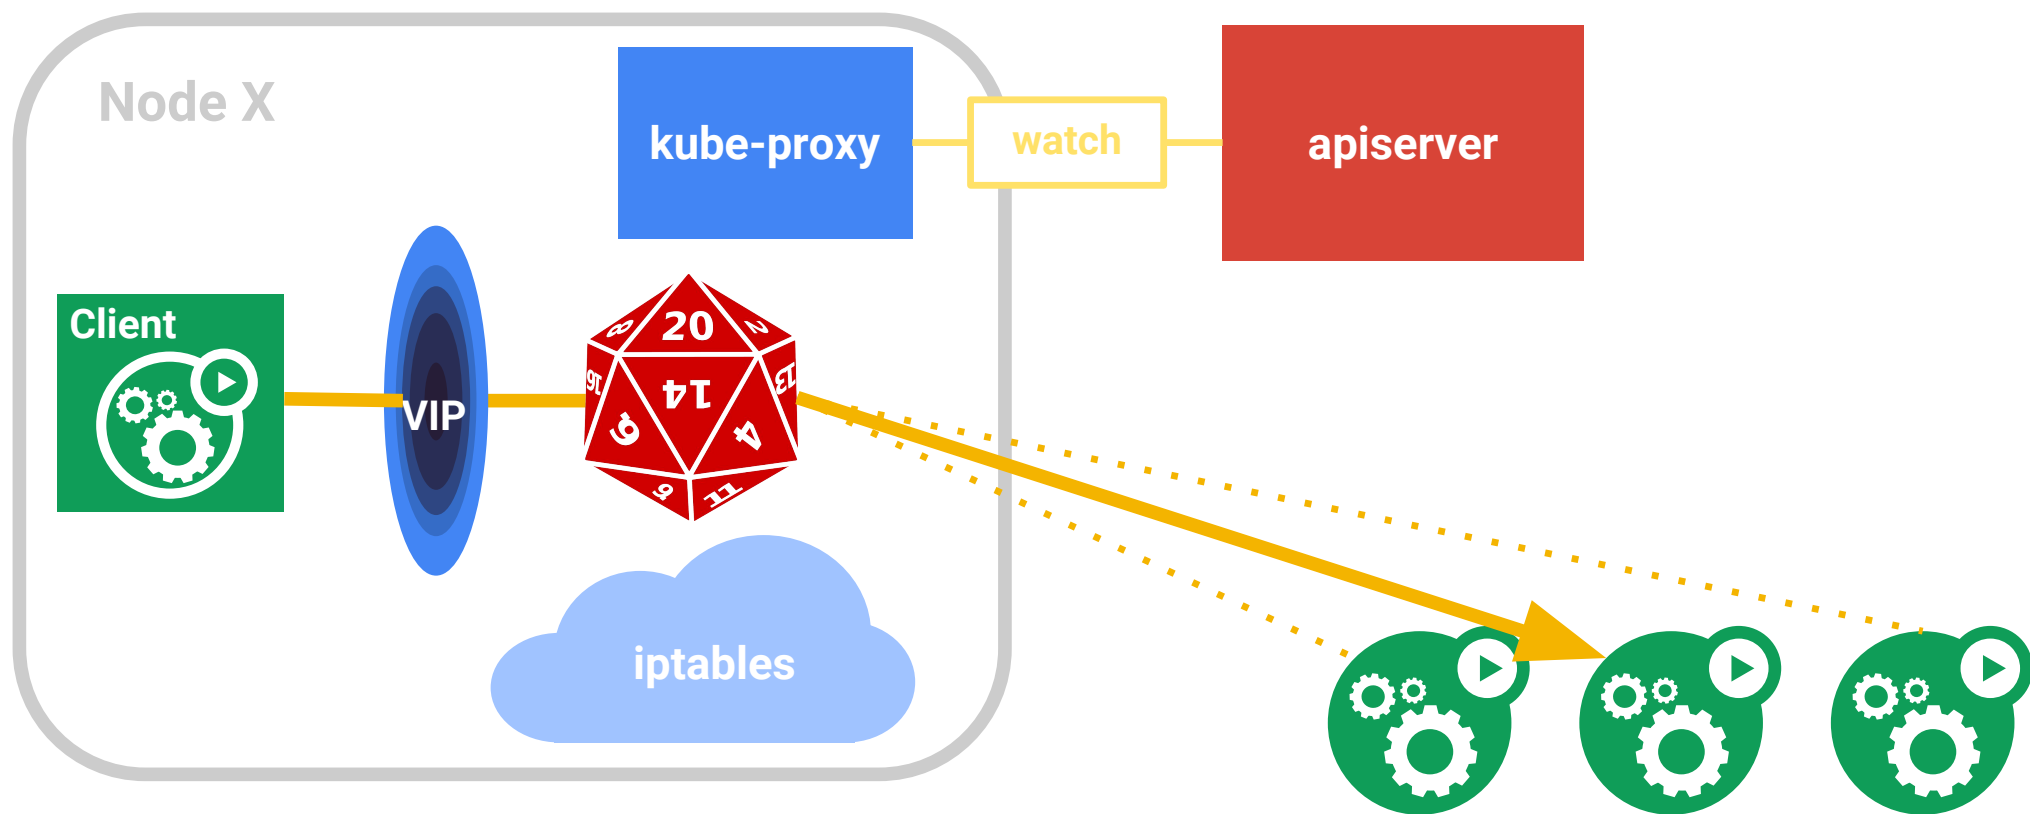
iptables kube-proxy



iptables kube-proxy



iptables kube-proxy



External services

Services VIPs are only available **inside** the cluster

Need to receive traffic from “the outside world”


Service “type”

- NodePort: expose on a port on every node
- LoadBalancer: provision a cloud load-balancer

DiY load-balancer solutions

- socat (for nodePort remapping)
- haproxy
- nginx

Ingress (L7 LB)



Ingress (L7 LB)

Many apps are HTTP/HTTPS

Services are L4 (IP + port)


Ingress maps incoming traffic to backend services

- by HTTP host headers
- by HTTP URL paths

HAProxy, NGINX, AWS and GCE implementations in progress

Now with SSL!

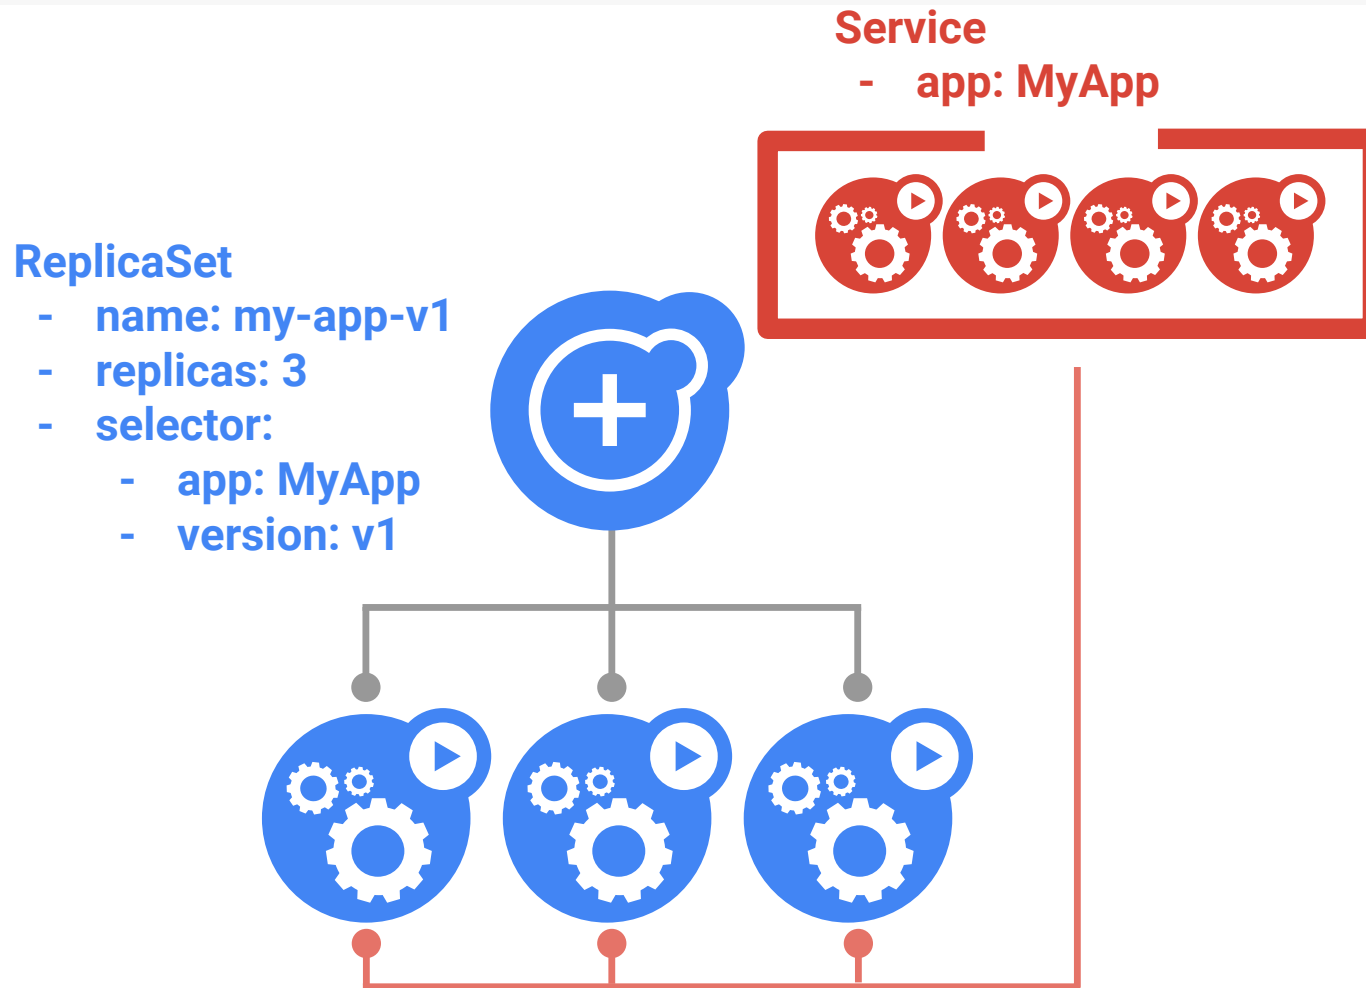
Status: BETA in Kubernetes v1.2



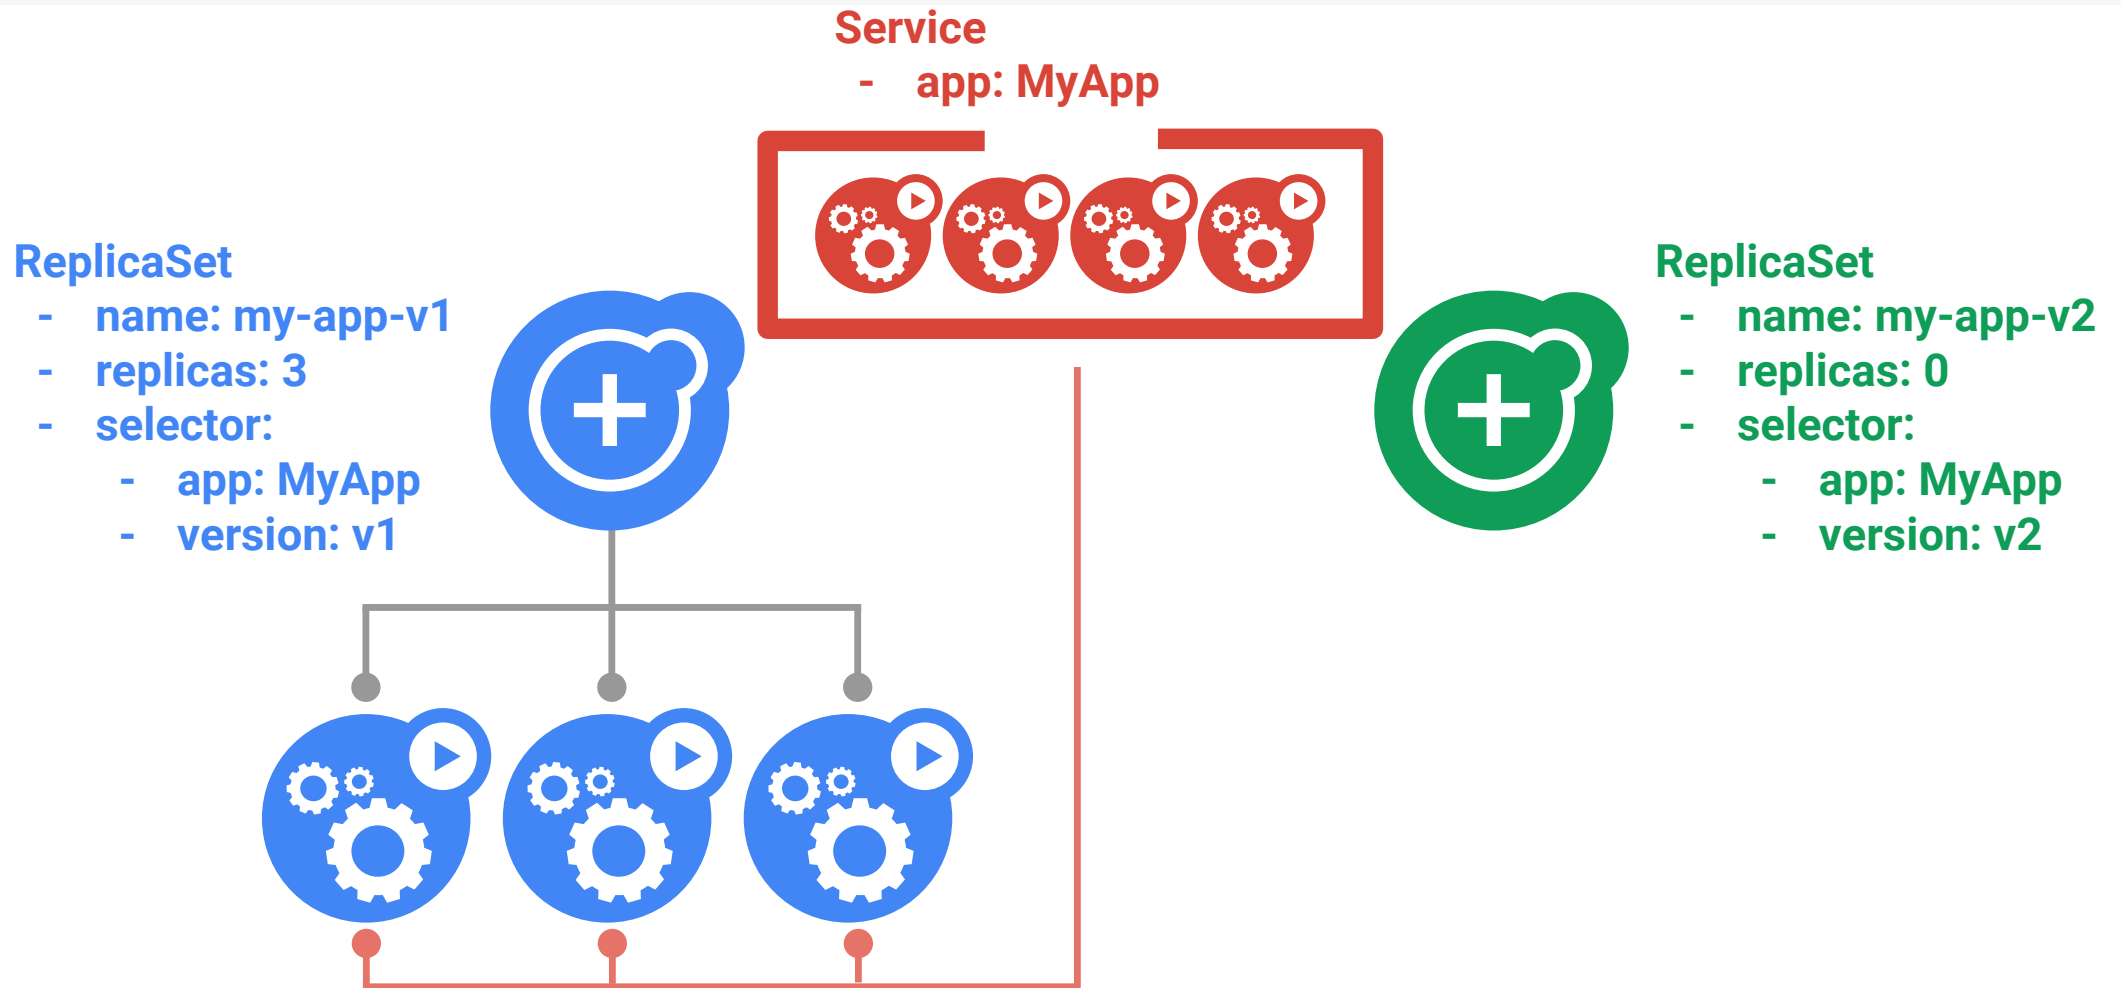
Rolling Update



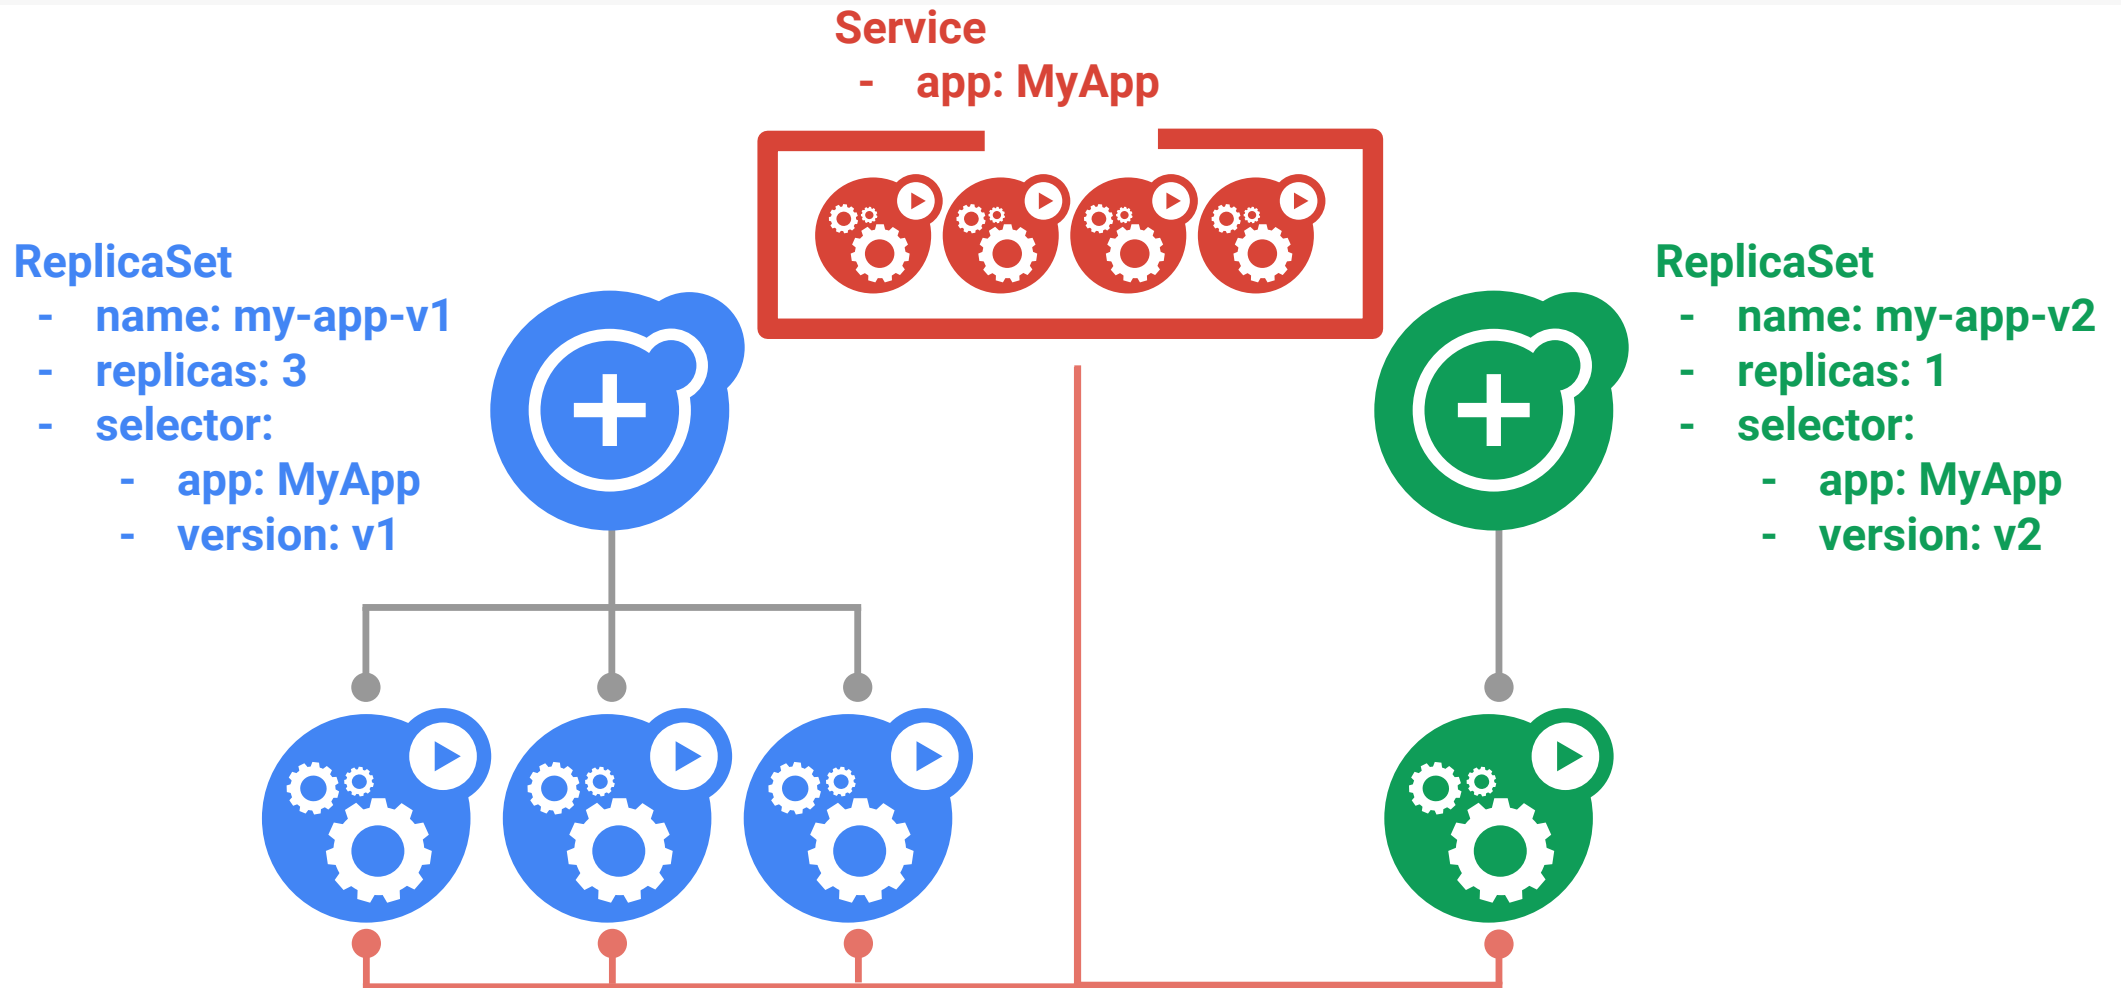
Rolling Update



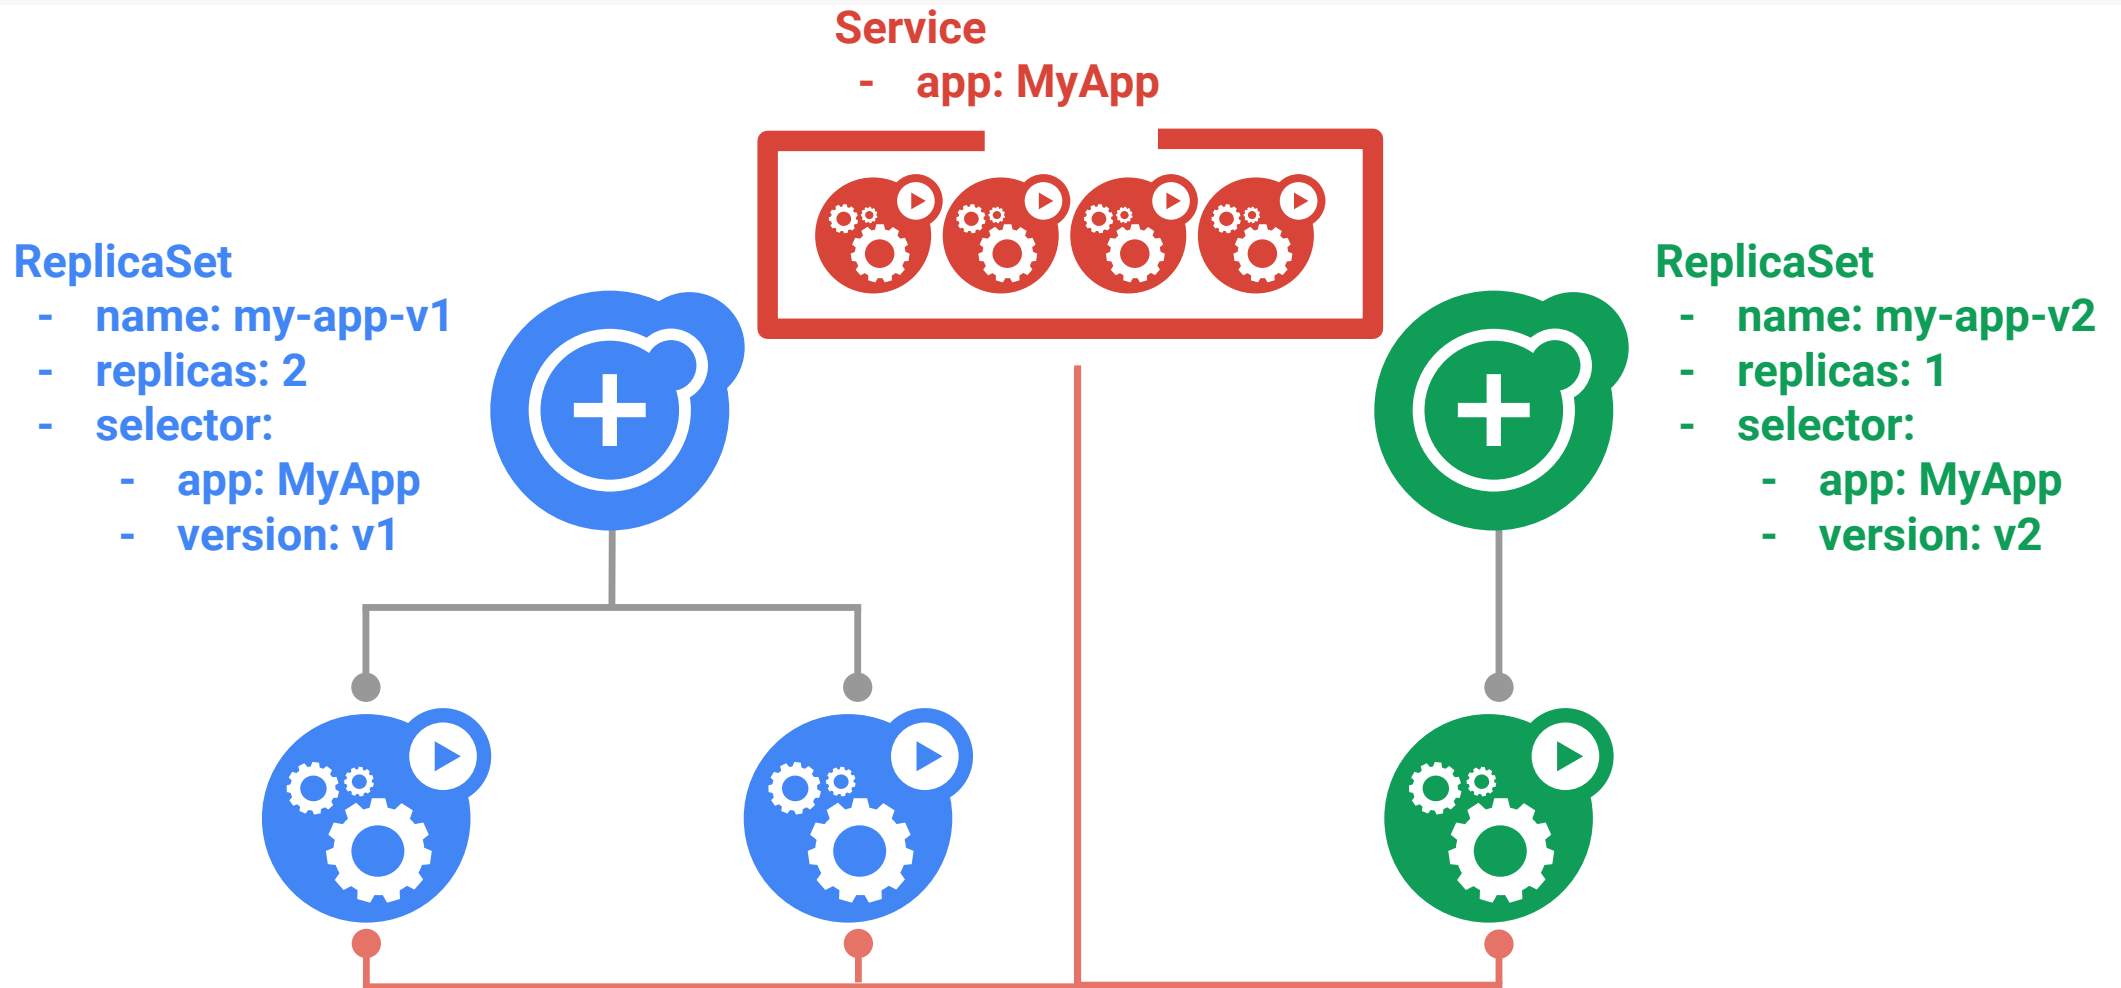
Rolling Update




Rolling Update



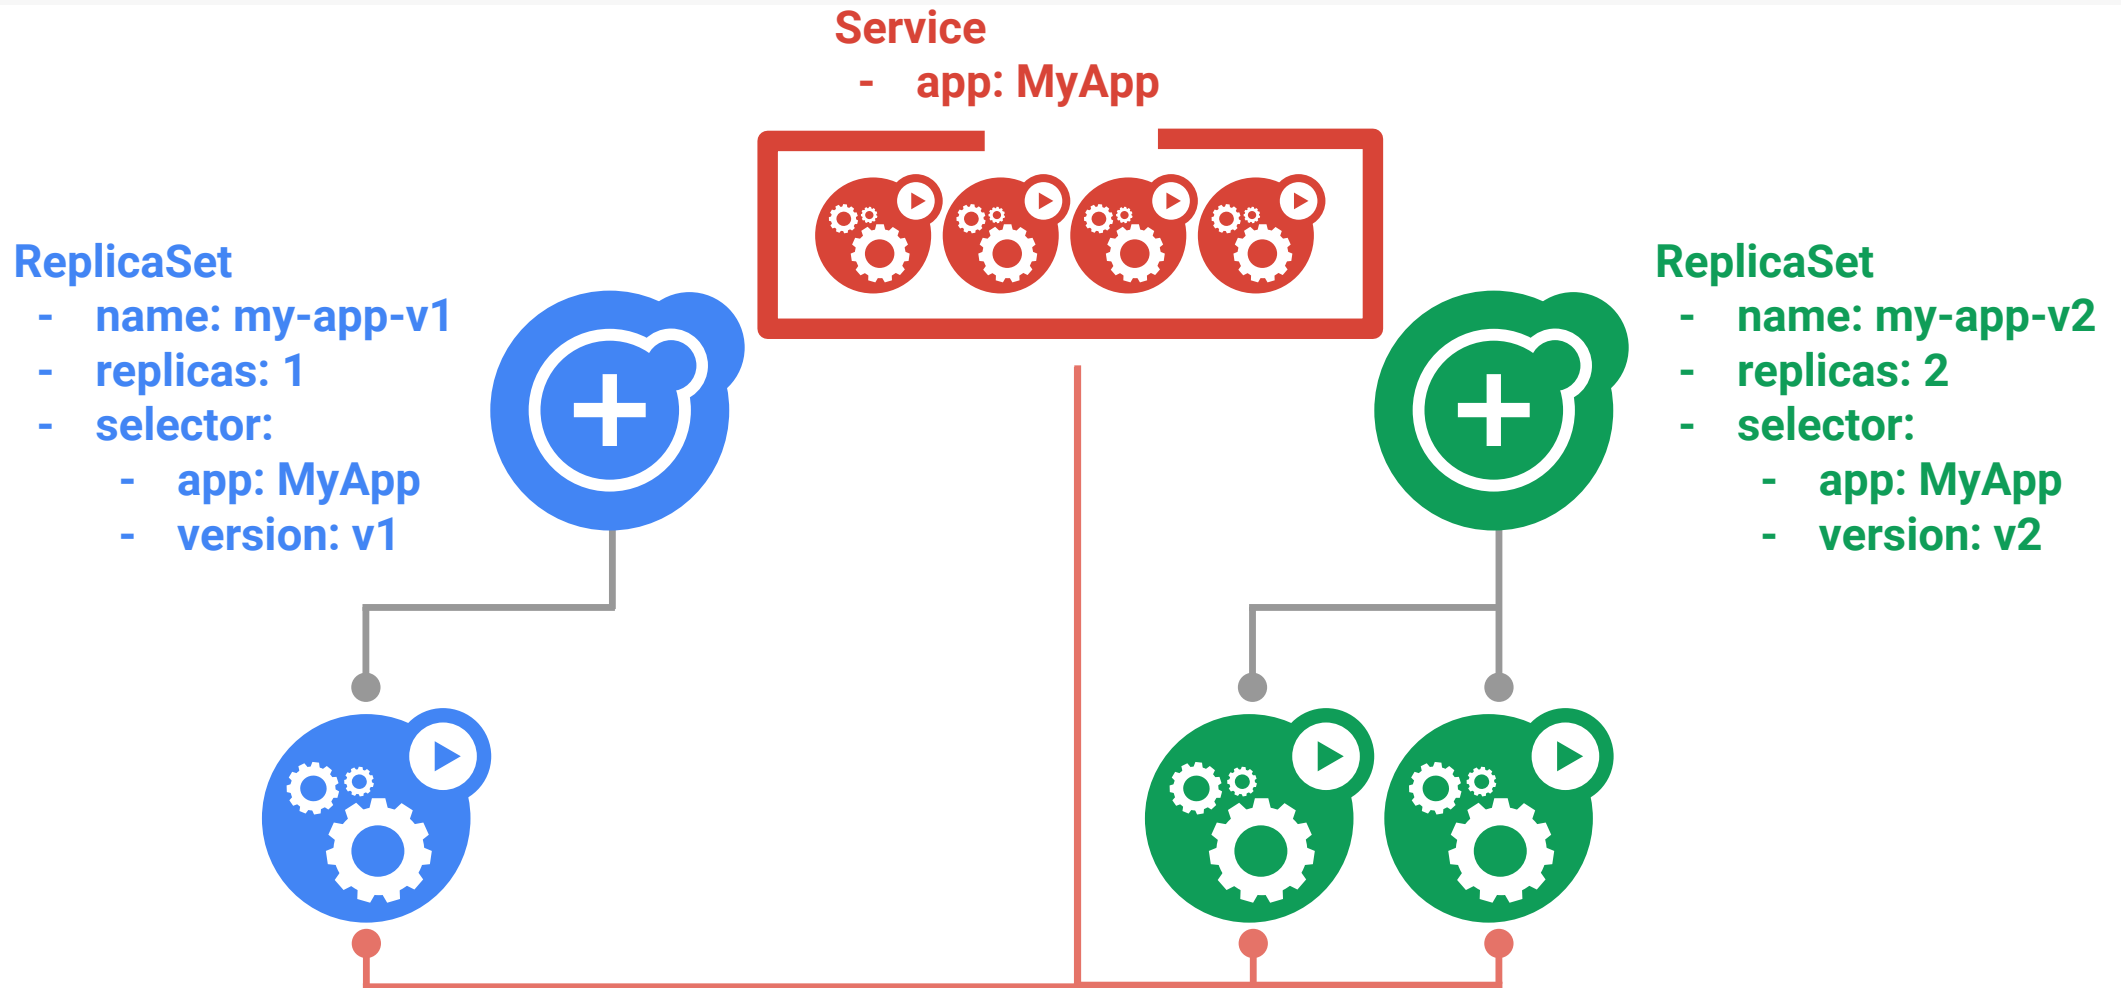
Rolling Update



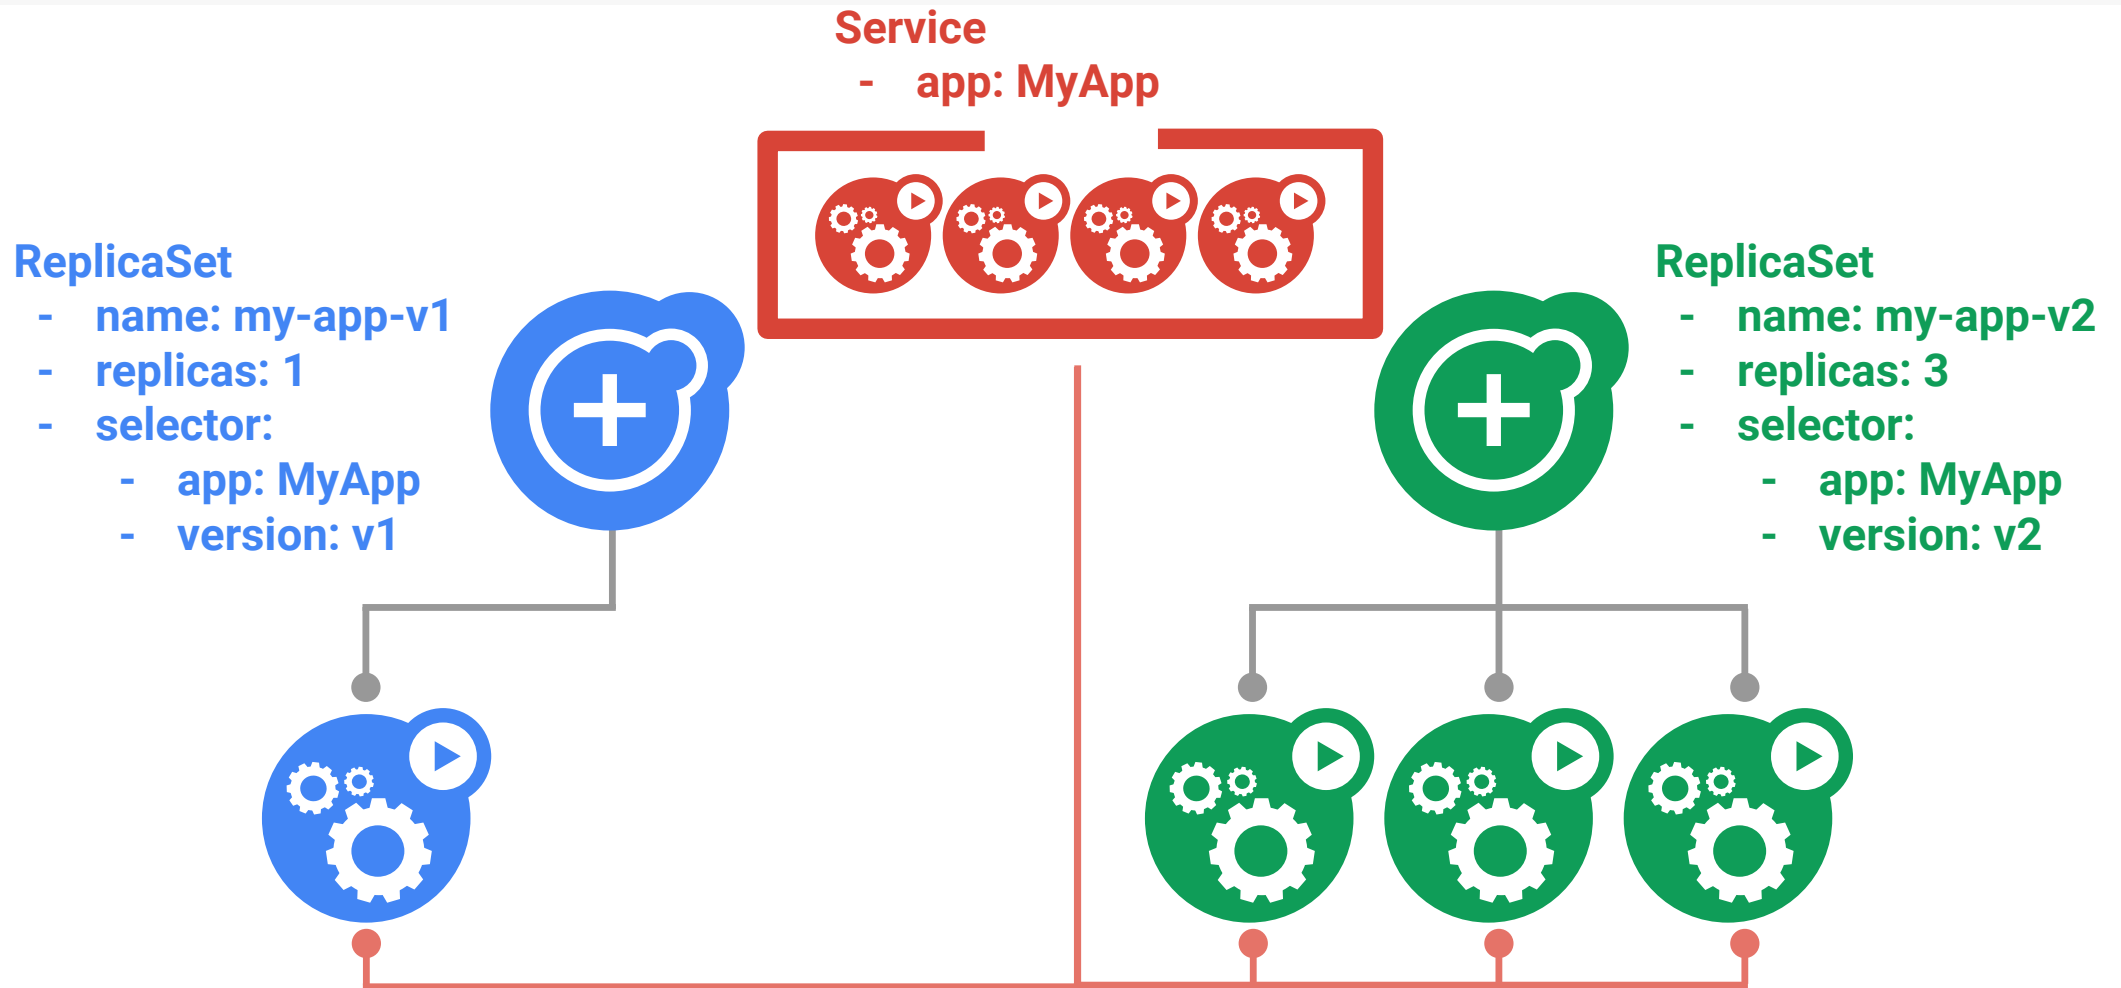
Rolling Update



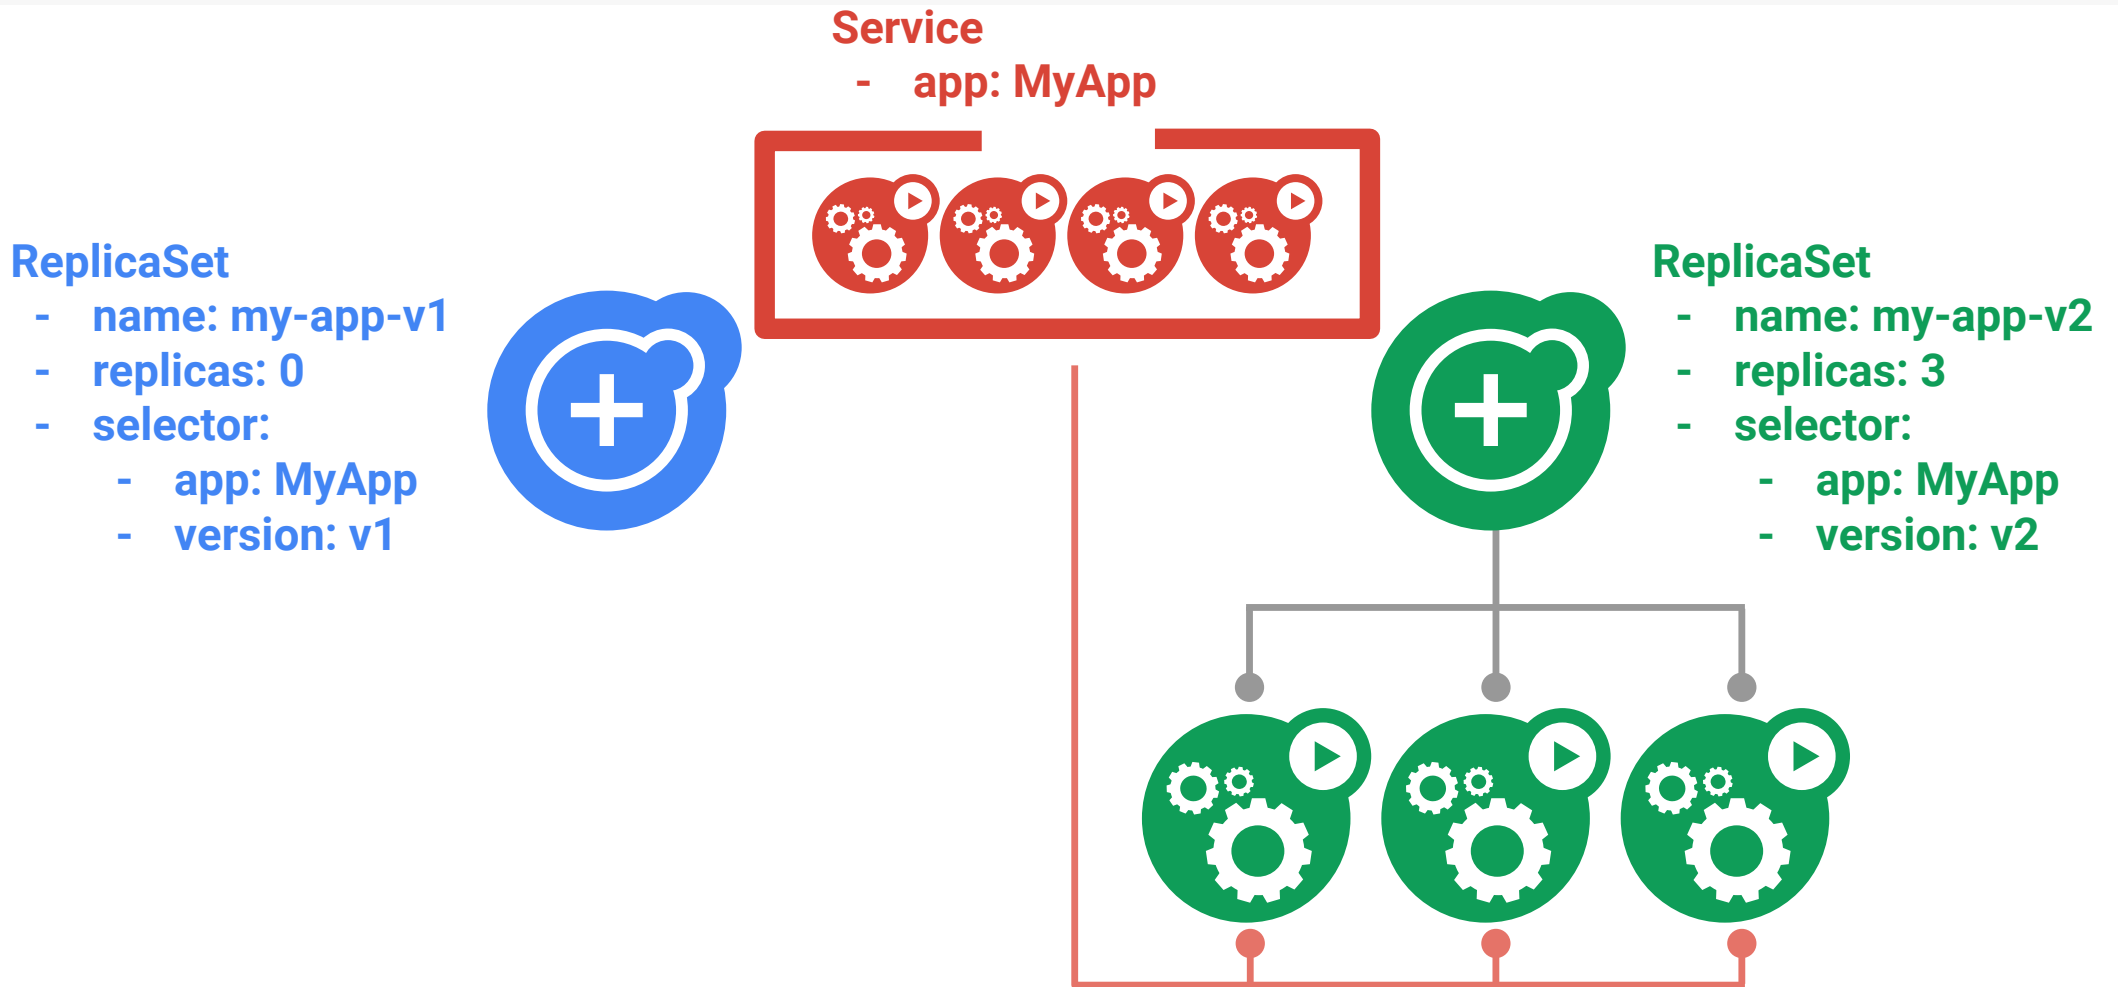
Rolling Update



Rolling Update



Rolling Update



Rolling Update


Service

- app: MyApp



ReplicaSet

- name: my-app-v2
- replicas: 3
- selector:
 - app: MyApp
 - version: v2



Deployments



Deployments

Updates-as-a-service

- Rolling update is imperative, client-side


Deployment manages replica changes for you

- stable object name
- updates are configurable, done server-side
- `kubectl edit` or `kubectl apply`

Aggregates stats

Can have multiple updates in flight

Status: BETA in Kubernetes v1.2



DaemonSets



DaemonSets

Problem: how to run a Pod on every node?

- or a subset of nodes


Similar to ReplicaSet

- principle: do one thing, don't overload

“Which nodes?” is a selector

Use familiar tools and patterns

Status: BETA in Kubernetes v1.2



Jobs



Jobs


Run-to-completion, as opposed to run-forever

- Express parallelism vs. required completions
- Workflow: restart on failure
- Build/test: don't restart on failure

Aggregates success/failure counts

Built for batch and big-data work

Status: GA in Kubernetes v1.2



PersistentVolumes



PersistentVolumes

A higher-level storage abstraction

- insulation from any one cloud environment


Admin provisions them, users claim them

- **NEW: auto-provisioning (alpha in v1.2)**

Independent lifetime from consumers

- lives until user is done with it
- can be handed-off between pods

Dynamically “scheduled” and managed, like nodes and pods



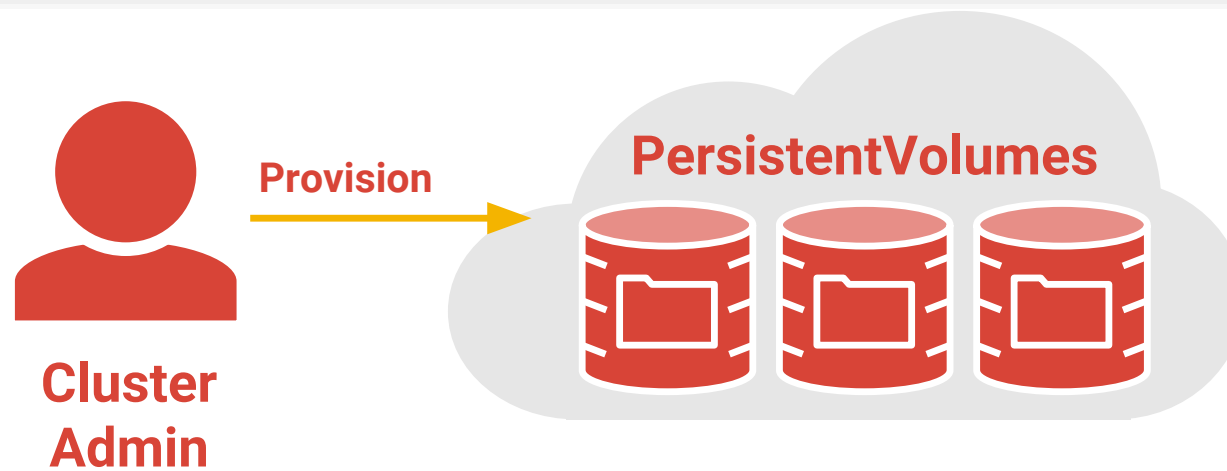
PersistentVolumes



**Cluster
Admin**



PersistentVolumes



PersistentVolumes



**Cluster
Admin**



User



PersistentVolumes



**Cluster
Admin**



User

Create



PVClaim


PersistentVolumes



**Cluster
Admin**



User



PersistentVolumes



Cluster Admin



PersistentVolumes



PVClaim



User

Create



Pod

Google Cloud Platform



PersistentVolumes



**Cluster
Admin**



User



PVClaim



Pod

Google Cloud Platform



PersistentVolumes



Cluster Admin



PVClaim



User

Delete



Pod

Google Cloud Platform




PersistentVolumes



**Cluster
Admin**



User



PersistentVolumes




Cluster Admin




User

Create



PVClaim



Pod



PersistentVolumes



Google Cloud Platform



PersistentVolumes



**Cluster
Admin**



User



PVClaim



Pod

Google Cloud Platform



PersistentVolumes



Cluster Admin



PersistentVolumes



PVClaim



User

Delete

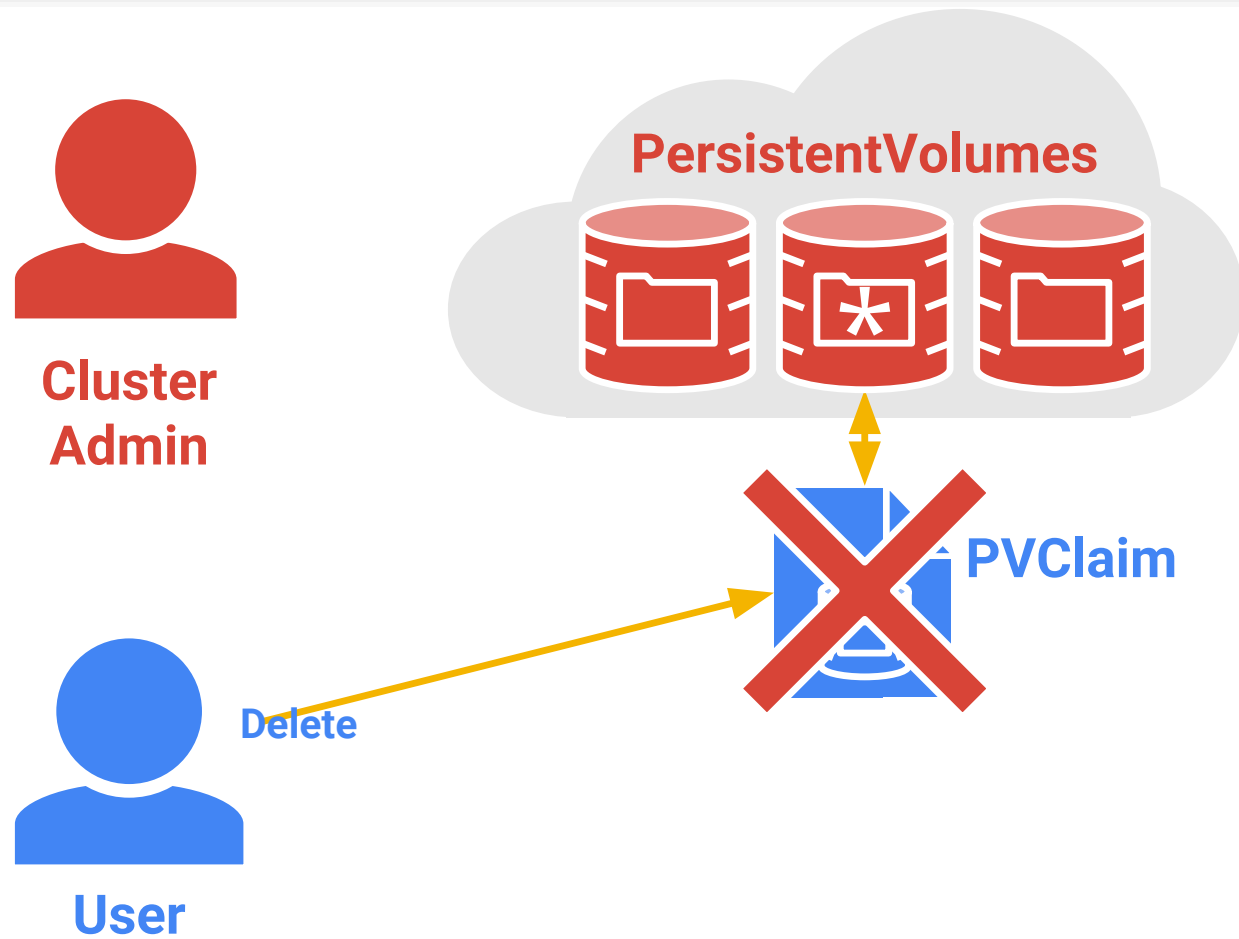


Pod

Google Cloud Platform



PersistentVolumes



PersistentVolumes



**Cluster
Admin**



User



Recycler



StatefulSets



StatefulSets


Goal: enable clustered software on Kubernetes

- mysql, redis, zookeeper, ...

Clustered apps need “identity” and sequencing guarantees

- stable hostname, available in DNS
- an ordinal index
- stable storage: linked to the ordinal & hostname
- discovery of peers for quorum
- startup/teardown ordering

Status: ALPHA in Kubernetes v1.3



ConfigMaps



ConfigMaps

Goal: manage app configuration

- ...without making overly-brittle container images

[12-factor](#) says config comes from the environment


- Kubernetes is the environment

Manage config via the Kubernetes API

Inject config as a virtual volume into your Pods

- late-binding, live-updated (atomic)
- also available as env vars

Status: GA in Kubernetes v1.2



Secrets



Secrets

Goal: grant a pod access to a secured something

- don't put secrets in the container image!


12-factor says config comes from the environment

- Kubernetes is the environment

Manage secrets via the Kubernetes API

Inject secrets as virtual volumes into your Pods

- late-binding, tmpfs - never touches disk
- also available as env vars



HorizontalPodAutoscalers



HorizontalPodAutoScalers

Goal: Automatically scale pods as needed


- based on CPU utilization (for now)
- custom metrics in Alpha

Efficiency now, capacity when you need it

Operates within user-defined min/max bounds

Set it and forget it

Status: GA in Kubernetes v1.2



Multi-Zone Clusters



Multi-Zone Clusters

Goal: zone-fault tolerance for applications


Zero API changes relative to kubernetes

- Create services, ReplicaSets, etc. exactly as usual

Nodes and PersistentVolumes are labelled with their availability zone

- Fully automatic for GKE, GCE, AWS
- Manual for on-premise and other cloud providers (for now)

Status: GA in Kubernetes v1.2



Namespaces



Namespaces

Problem: I have too much stuff!

- name collisions in the API
- poor isolation between users
- don't want to expose things like Secrets

Solution: Slice up the cluster

- create new Namespaces as needed
 - per-user, per-app, per-department, etc.
- part of the API - NOT private machines
- most API objects are namespaced
 - part of the REST URL path
- Namespaces are just another API object
- One-step cleanup - delete the Namespace
- Obvious hook for policy enforcement (e.g. quota)




Resource Isolation



Resource Isolation

Principles:

- Apps must not be able to affect each other's performance
 - if so it is an **isolation failure**
- Repeated runs of the same app should see ~equal behavior
- QoS levels drives resource decisions in (soft) real-time
- Correct in all cases, optimal in some
 - reduce unreliable components
- SLOs are the lingua franca



Strong isolation

Pros:

- Sharing - users don't worry about interference (aka the noisy neighbor problem)
- Predictable - allows us to offer strong SLAs to apps

Cons:

- Stranding - arbitrary slices mean some resources get lost
- Confusing - how do I know how much I need?
 - analog: what size VM should I use?
 - smart auto-scaling is needed!
- Expensive - you pay for certainty

In reality this is a multi-dimensional bin-packing problem: CPU, memory, disk space, IO bandwidth, network bandwidth, ...



Requests and Limits

Request:

- how much of a resource you are asking to use, with a strong guarantee of availability
 - CPU (seconds/second)
 - RAM (bytes)
- scheduler will not over-commit requests

Limit:

- max amount of a resource you can access

Repercussions:

- Usage > Request: resources **might** be available
- Usage > Limit: throttled or killed



Quality of Service

Defined in terms of Request and Limit

Guaranteed: highest protection

- `request > 0 && limit == request`

Burstable: medium protection


- `request > 0 && limit > request`

Best Effort: lowest protection

- `request == 0`

What does “protection” mean?

- OOM score
- CPU scheduling



Quota and Limits



ResourceQuota

Admission control: apply limits in **aggregate**

Per-namespace: ensure no user/app/department abuses the cluster

Reminiscent of disk quota by design

Applies to each type of resource

- CPU and memory for now

Disallows pods without resources



LimitRange


Admission control: limit the limits

- min and max
- ratio of limit/request

Default values for unspecified limits

Per-namespace

Together with ResourceQuota gives cluster admins powerful tools



Cluster Auto-Scaling



Cluster Autoscaler


Add nodes when needed

- there are pending pods
- some pending pods would fit if we add a node

Remove nodes when not needed

- after removal, all pods must fit remaining nodes

Status: Works on GCE, GKE and AWS



Scalability




Scalability & Performance

SLO met at <2000 nodes, <60000 pods

- 99% of API calls return in < 1 second
- 99% of pods start in < 5 seconds

Coming soon

- protobufs in API storage (already enabled on the wire)
- 5000 nodes



Design principles

Declarative > imperative: State your desired results, let the system actuate

Control loops: Observe, rectify, repeat

Simple > Complex: Try to do as little as possible

Modularity: Components, interfaces, & plugins

Legacy compatible: Requiring apps to change is a non-starter

Network-centric: IP addresses are cheap

No grouping: Labels are the only groups

Sets > Pets: Manage your workload in bulk

Open > Closed: Open Source, standards, REST, JSON, etc.



Kubernetes (K8s) Community

~5k Commits
in 1.4 over 3
months

> 800 Unique
Contributors


Top 0.01% of
all Github
Projects


2500+ External
Projects Based
on K8s

Companies Contributing



Companies Using





“Niantic chose GKE for its ability to **orchestrate their container cluster at planetary-scale**, freeing its team to focus on deploying live changes for their players.” - Niantic



Further Reading

If this talk was interesting, deeper academic reading on cluster management:

“Borg, Omega, and Kubernetes”

ACM Queue, March 2, 2016, [Volume 14, issue 1](#)

<http://queue.acm.org/detail.cfm?id=2898444>

Or a hands-on “Hello World” quickstart to build a Docker image and run it on a Kubernetes cluster:

<http://kubernetes.io/docs/hellonode/>

Another hard problem: how do you run N Kubernetes clusters as a service?

- create/delete, update, monitor, repair, escalate, upgrade, backup/restore, zonal isolation, incremental rollouts, support ticket escalation, provisioning, and more!



Questions?

Potential discussion:

- What about Docker Swarm?
- ... Mesos?
- What's next for Kubernetes and Container Engine?
- Why Google not FB/Uber/MS/Ama/etc?
- How do I get an internship / job?
 - Let's discuss!

- Alex on Philosophy:
 - Imperative vs. declarative
 - Orchestration vs. choreography
 - Product vs. tech
 - User guide vs. design doc
 - Engineering code vs. organizations
 - Your team is a design parameter
 - Launch and iterate; MVP

More questions?

Happy to chat!

- Lunch
- 1:1's after that
- mohr@google.com
- 590s@alexmohr.com

Google Cloud Platform